

# **Aprendizado Profundo (Deep Learning)**

**Similarity Learning**

# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks
- Improving Similarity Learning
- Change Detection

# Limitation of Classification Approaches

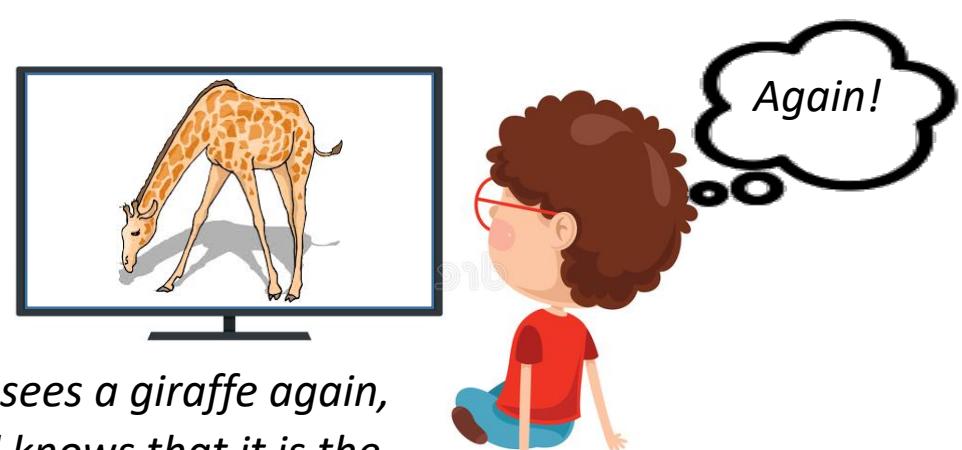
Typically, a classification involves fitting a model given many (some times, hundreds of thousands) examples of each class.

*Humans can do better!*



*A child knows that it is something unknown when he first sees a giraffe*

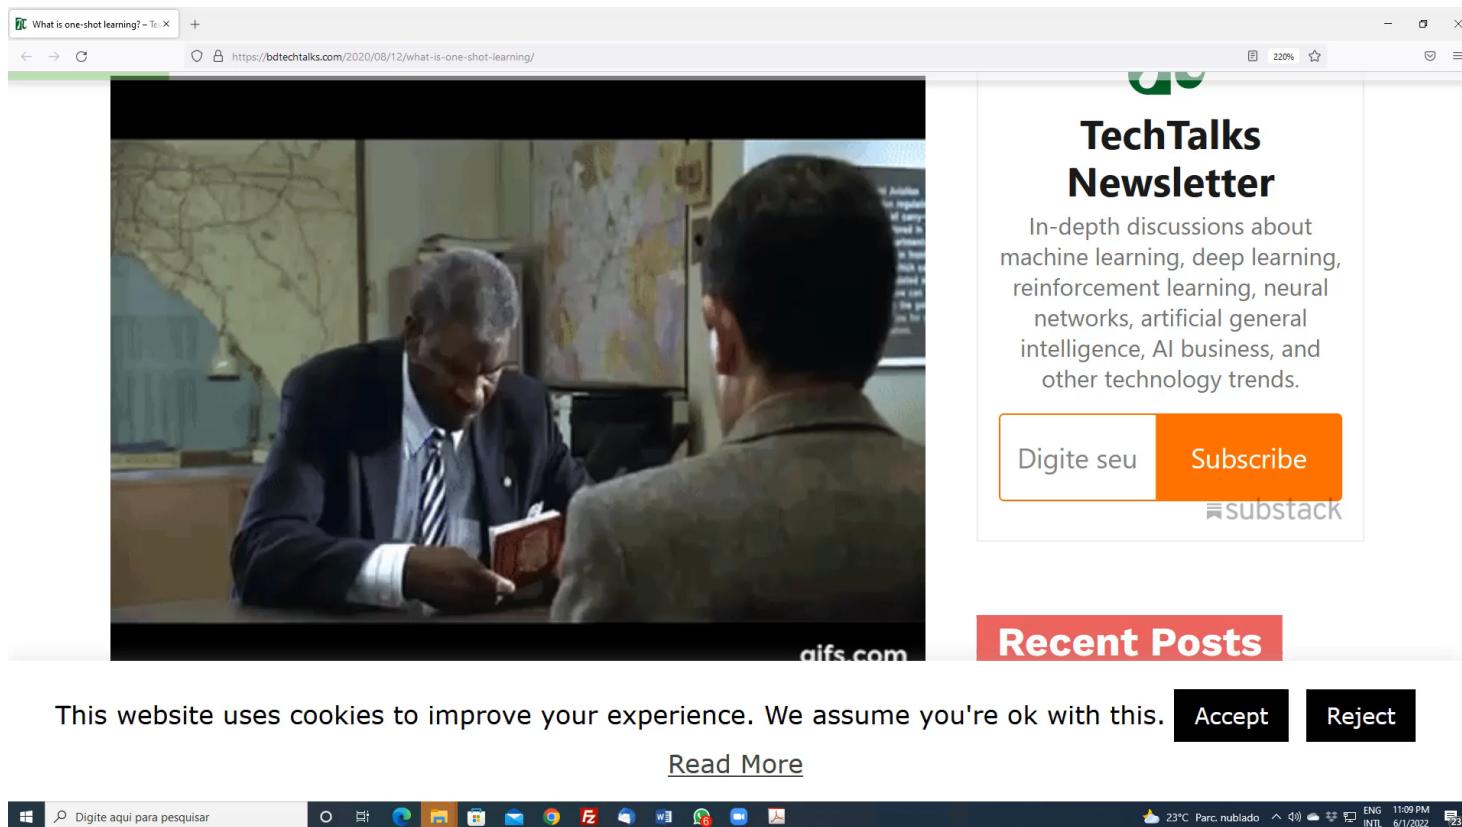
*Later on*



*when he sees a giraffe again, the child knows that it is the same animal he saw earlier.*

# One-Shot Learning

One-shot learning is about looking at two images never seen before and saying whether they represent the same object.



# One-Shot Learning applications

## Biometrics



# One-Shot Learning applications

## Recommendation systems



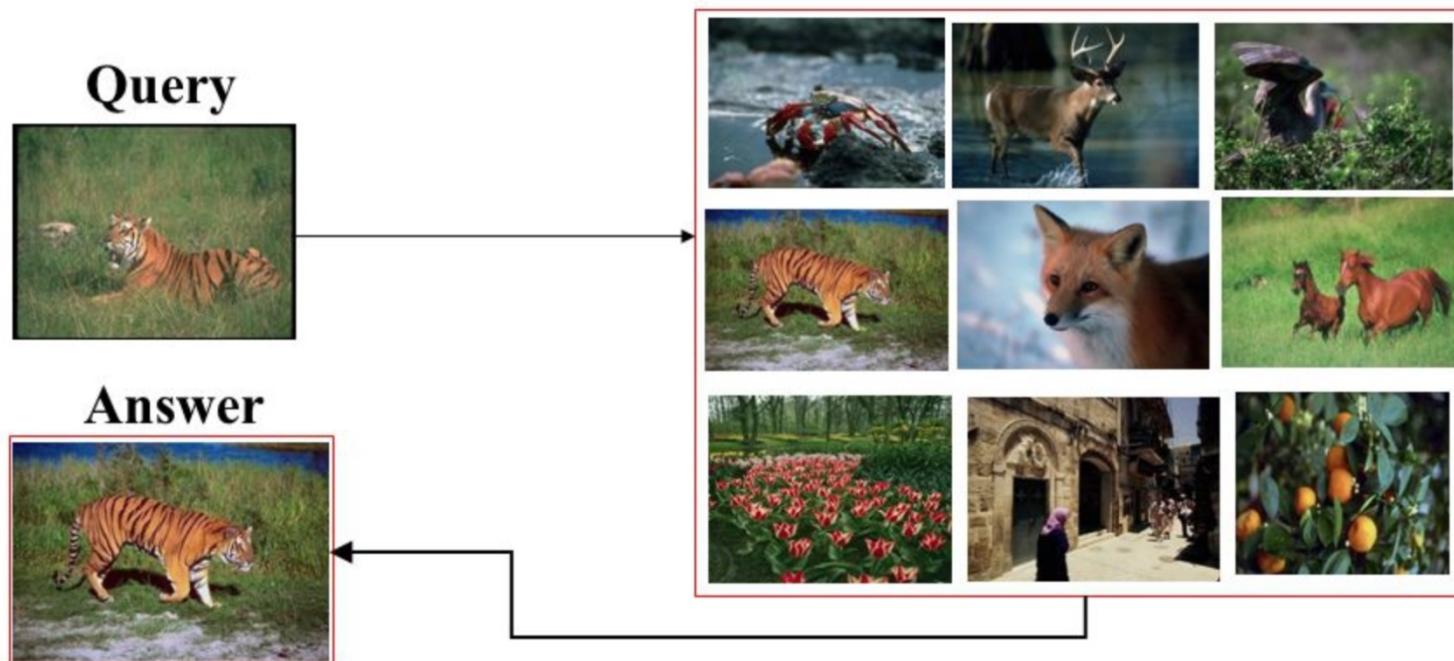
# One-Shot Learning applications

Matching résumés to job descriptions



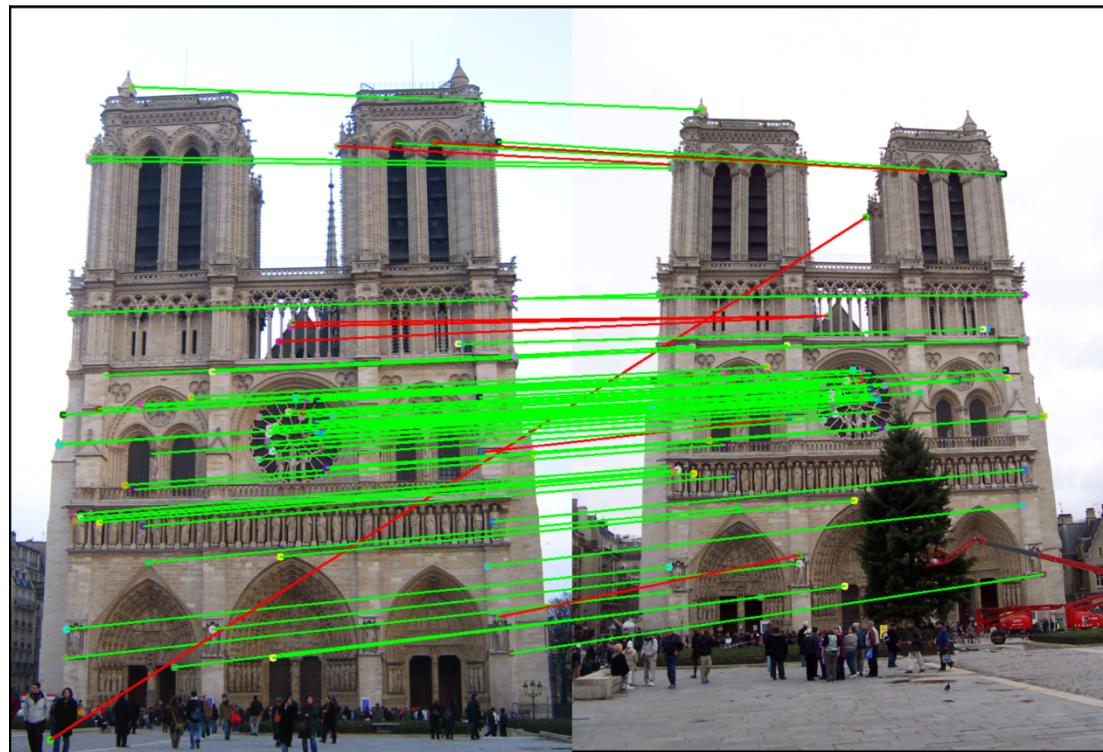
# One-Shot Learning applications

## Content-based image retrieval



# One-Shot Learning applications

## Image Matching



# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks
- Improving Similarity Learning
- Change Detection

# Face Recognition

**Identification problem:** who is the subject?

**A possible solution:**

- Define a class for each subject in the gallery.
- Train a classifier
- Use the trained classifier to answer the question.



**What if you add a new subject in the gallery?**

**Repeat all steps again!**

Gallery			

# Face Recognition

**Verification problem:** has the subject the alleged identity?

Are they the same person?

A



**NO**

B



C



**YES**

D

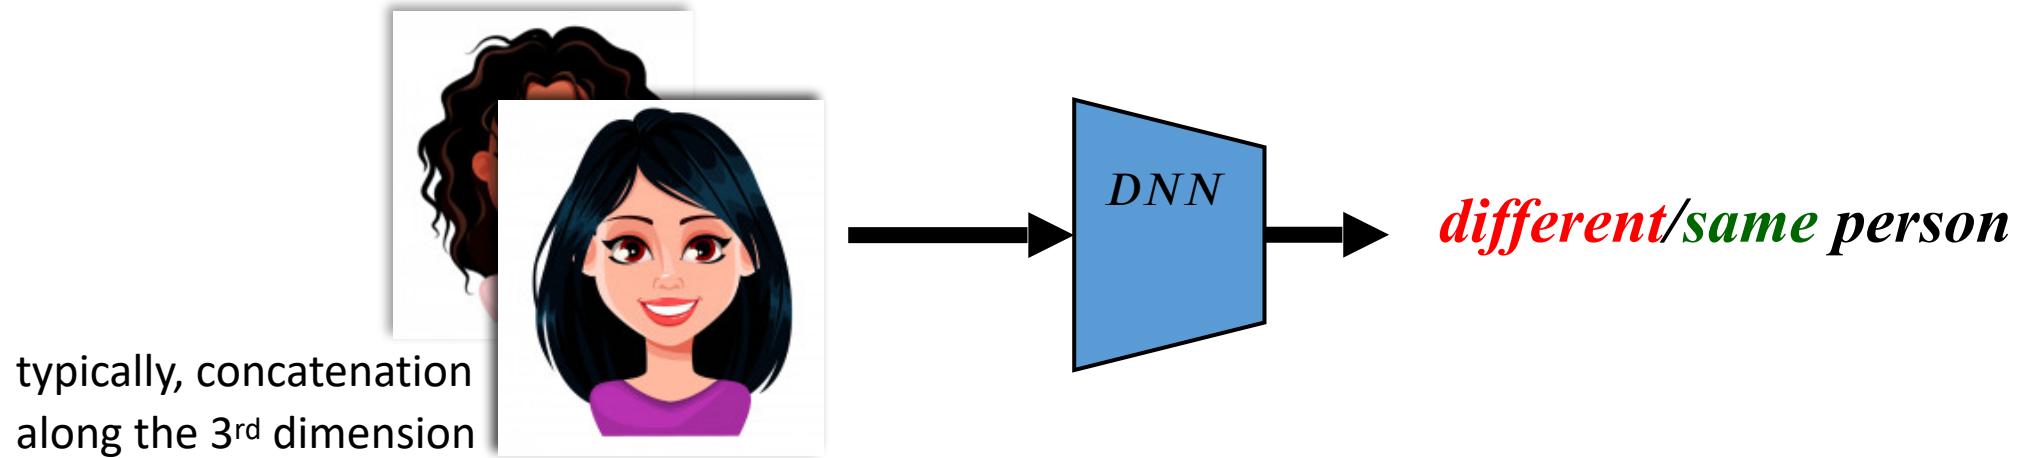


# Content

- One-shot Learning
- Example: face recognition
- **Similarity Learning**
- Siamese Networks
- Improving Similarity Learning
- Change Detection

# Early Fusion

A simple solution for verification: transform into a binary classification problem



use the binary **cross-entropy** as loss function.

# Similarity Learning

Instead of learning how to classify the input – **learn how to measure similarity** of two inputs

Are they the same person?



**Low** similarity score



**High** similarity score



# Similarity Learning

## Learning Distance Metric

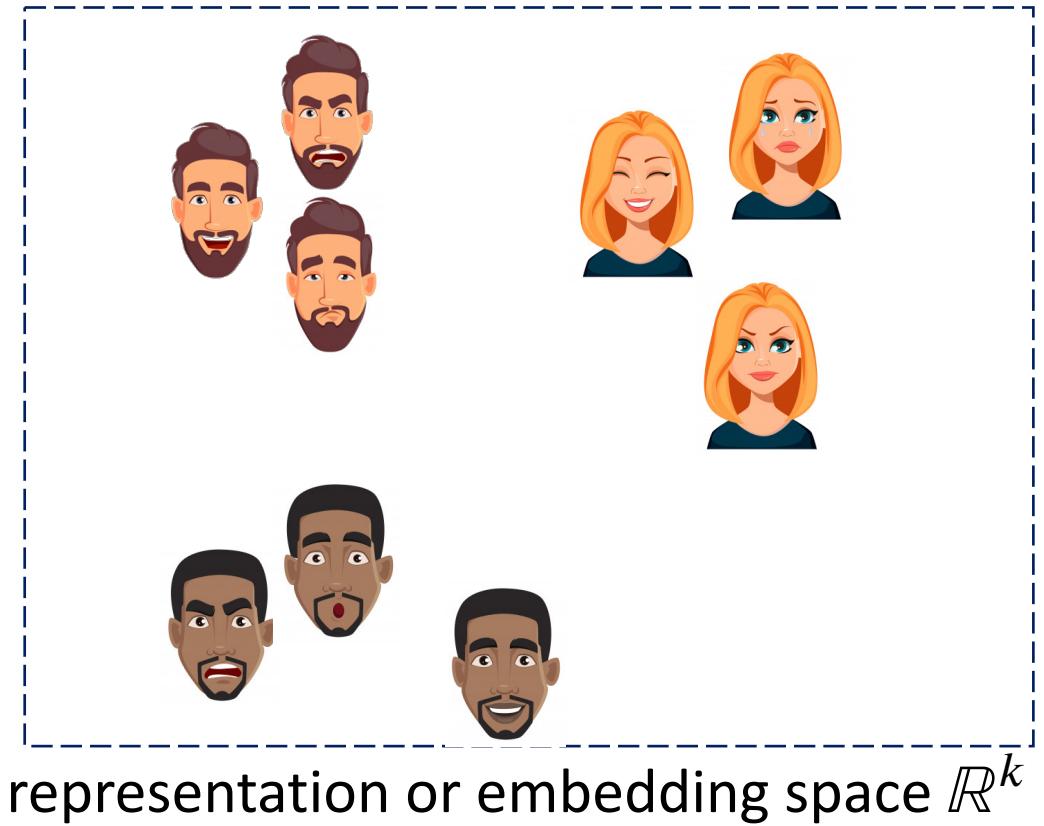
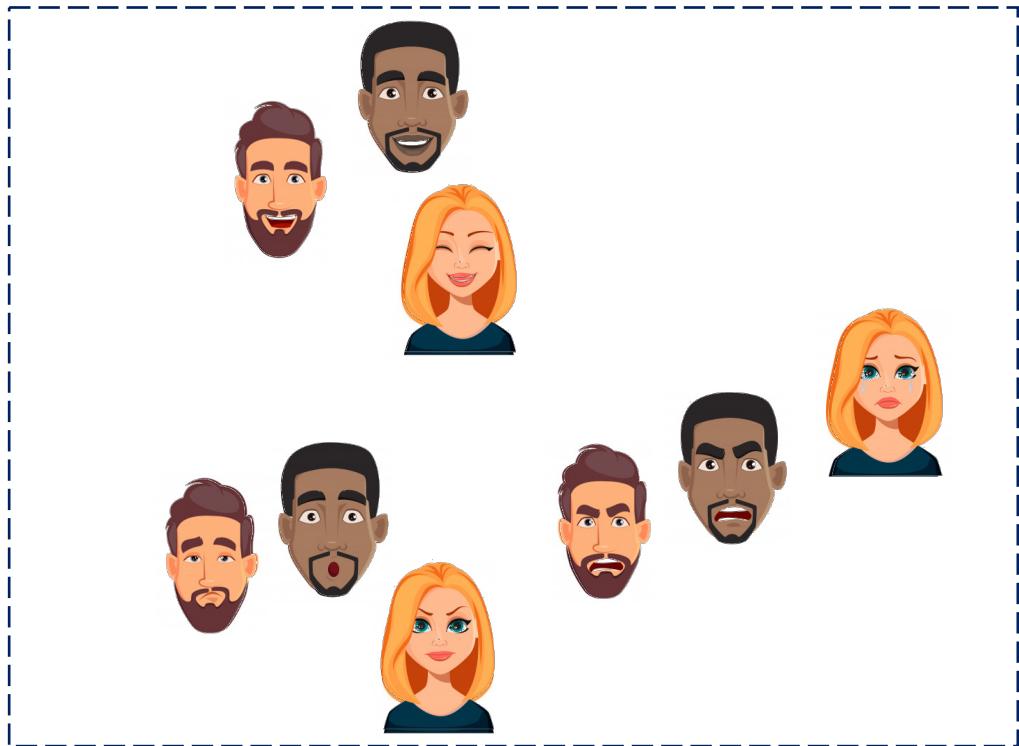
- An active research topic in CV community.
- Learn a **distance function** that assigns small/large distance to pairs of examples that are semantically similar/dissimilar.

## Learning Feature Embeddings

- With the advent of DNNs, metric learning turned to learning **feature embeddings** that better fits a simple distance function, such as Euclidean distance or cosine distance.

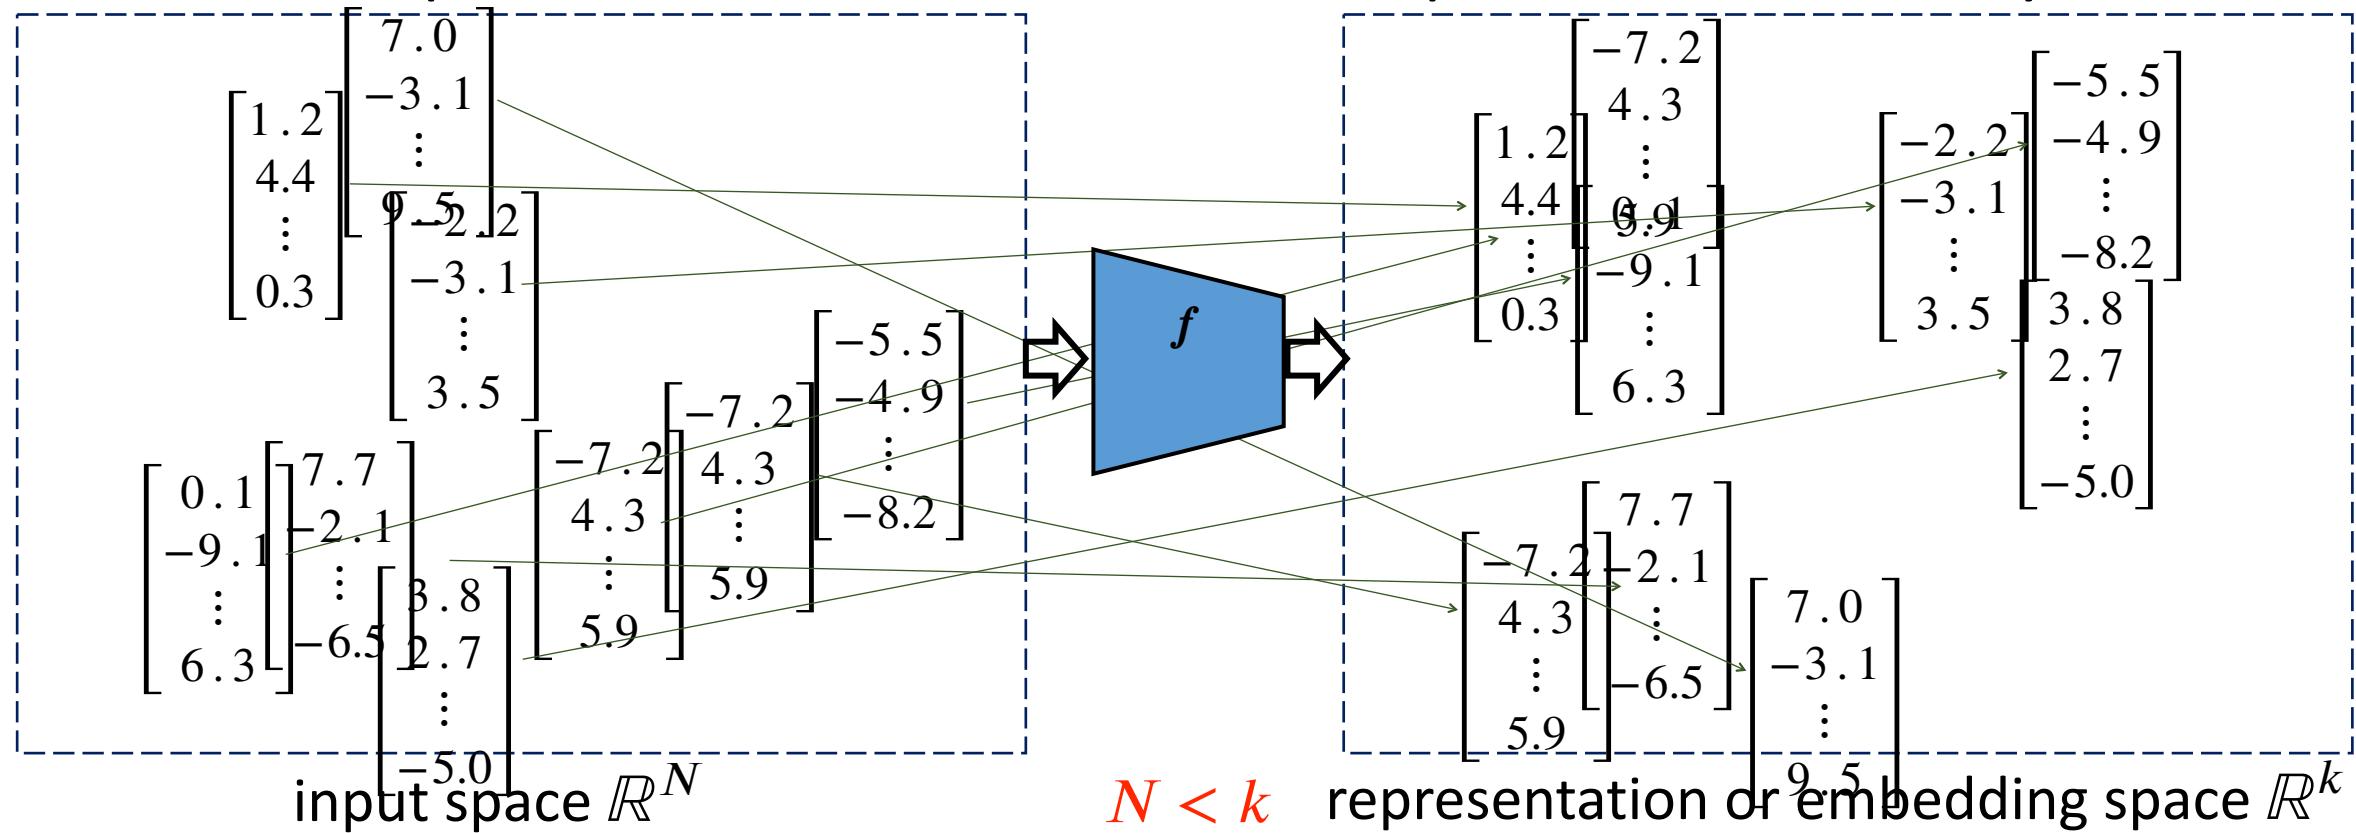
# Similarity Learning

Is about learning an **embedding space**  $\mathbb{R}^k$  where images of the same person are close and of different persons are far away.



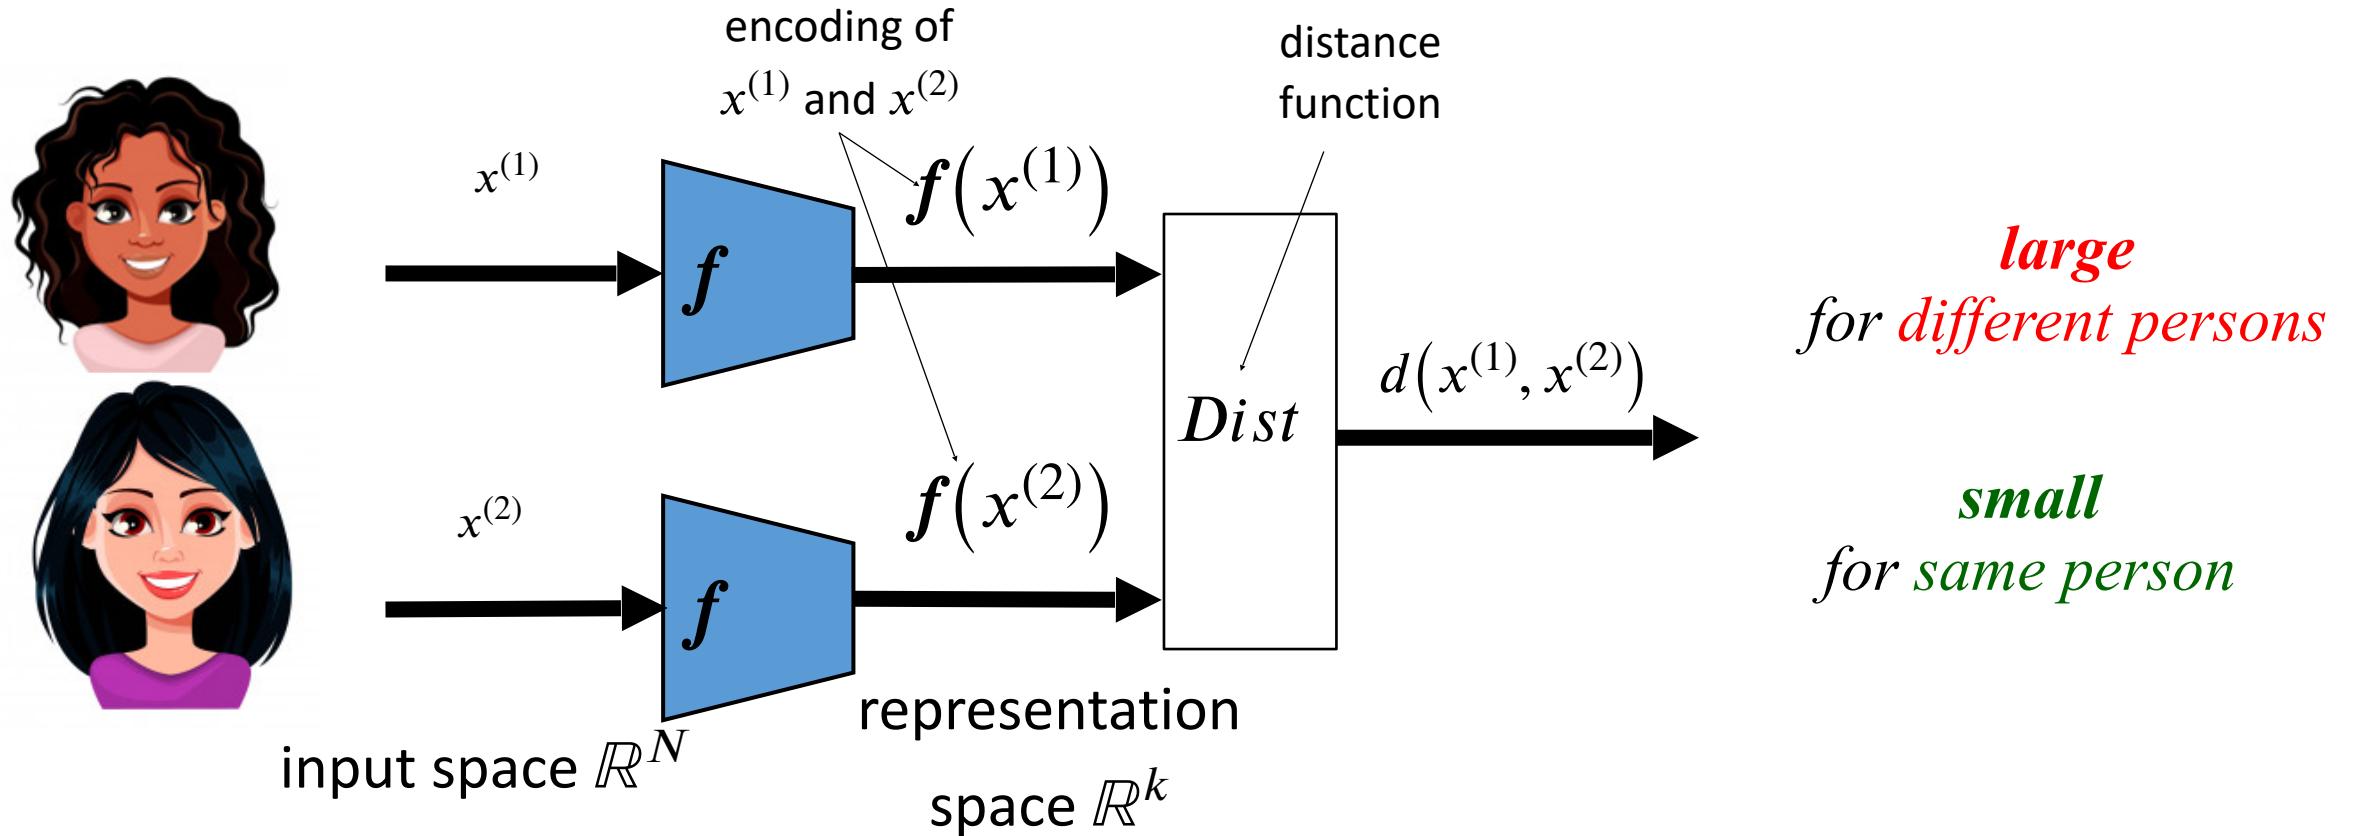
# Learning an Embedding Space

**Step 1:** find a mapping  $f: \mathbb{R}^N \rightarrow \mathbb{R}^k$  to an **embedding space** where images of the same person are close and of different persons are far away.



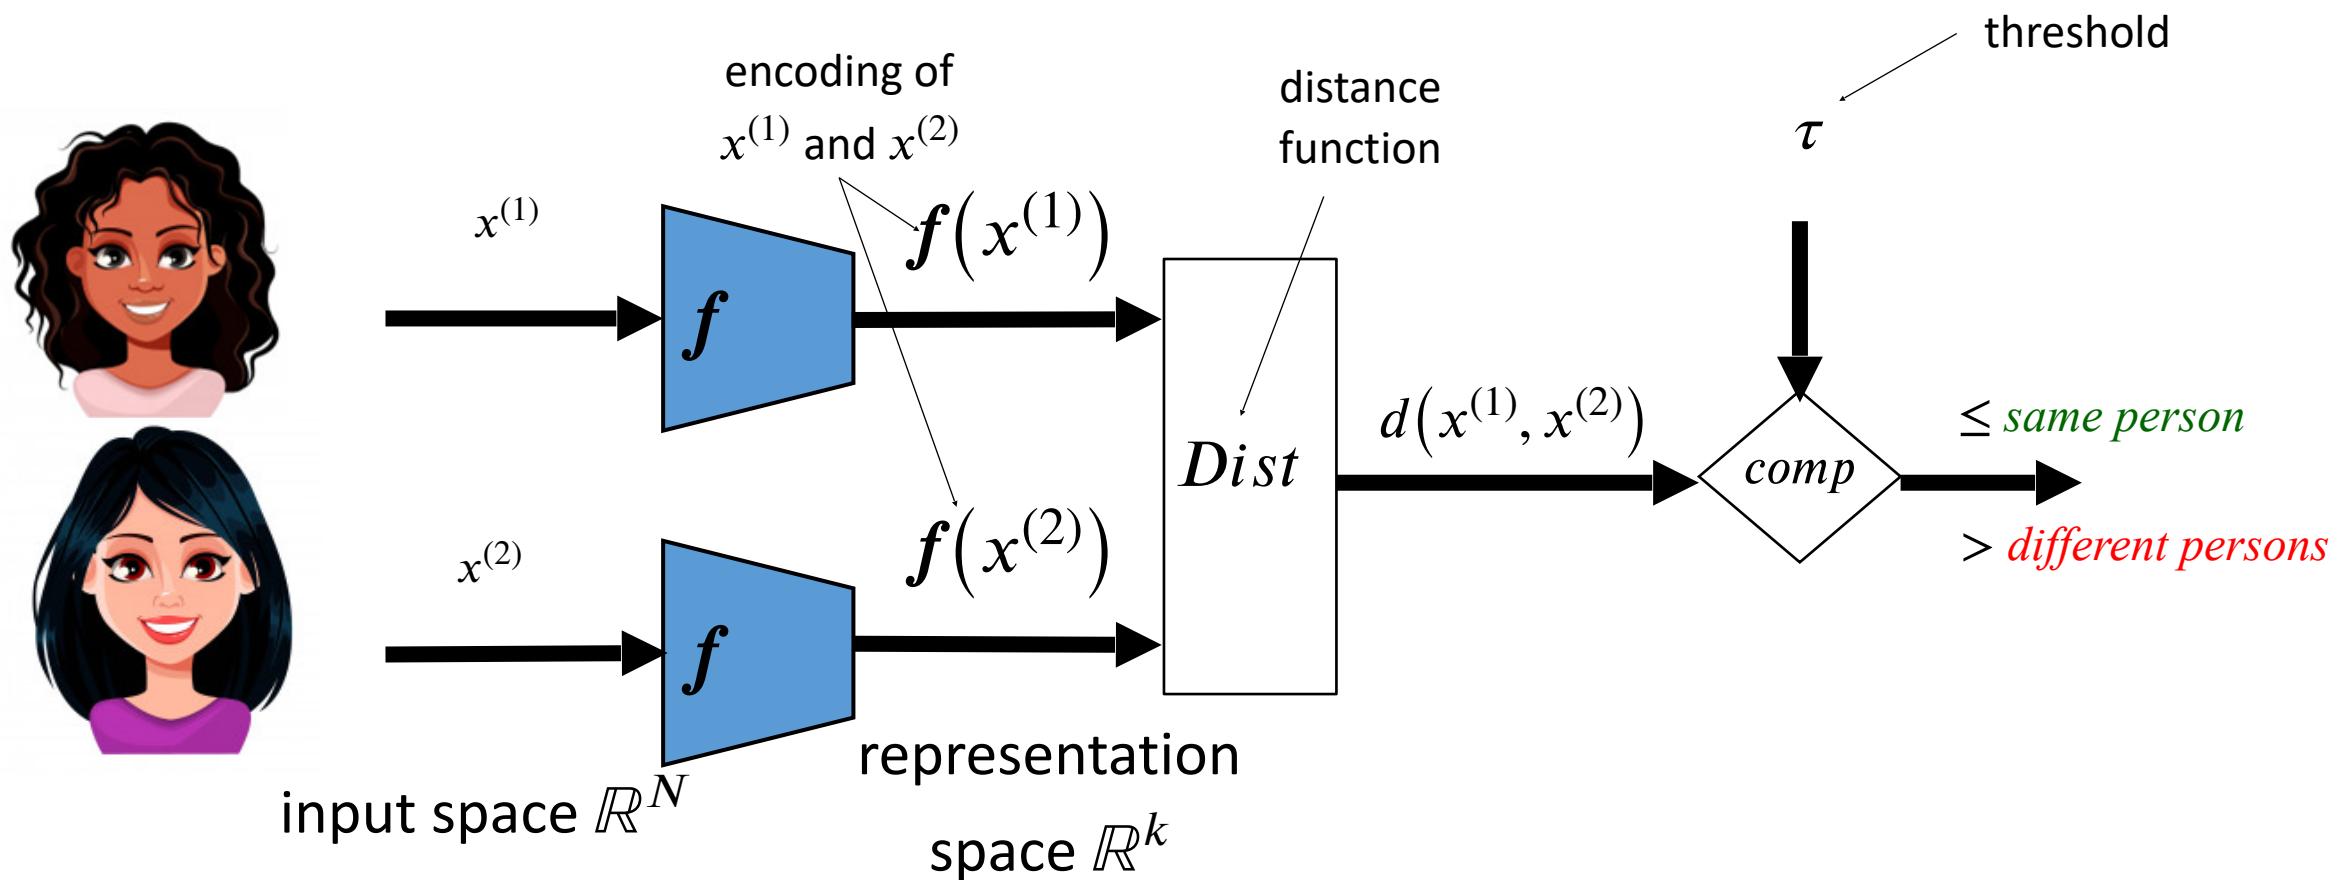
# Distance Learning

**Step 2:** choose a distance function  $Dist$  in the embedding space related to the (dis)similarity between input images  $x^{(1)}$  and  $x^{(2)}$ .



# Choose a Threshold

**Step 3:** define a threshold  $\tau$  to decide whether the inputs are of the same or of different persons.



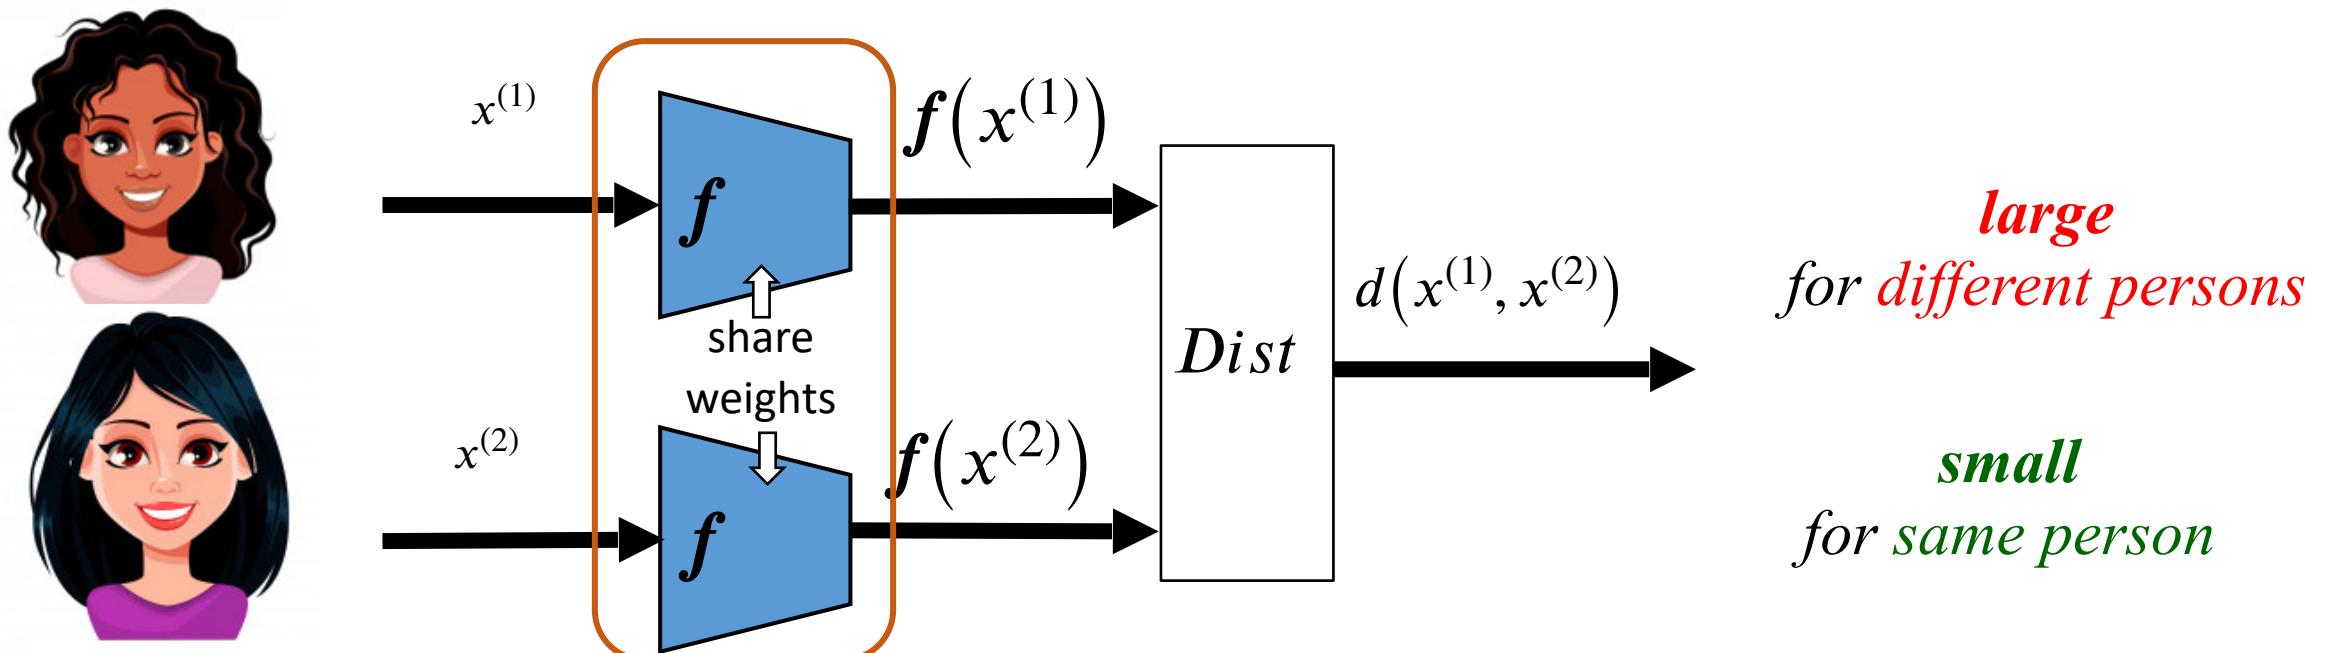
# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks
- Improving Similarity Learning
- Change Detection

# Siamese Network

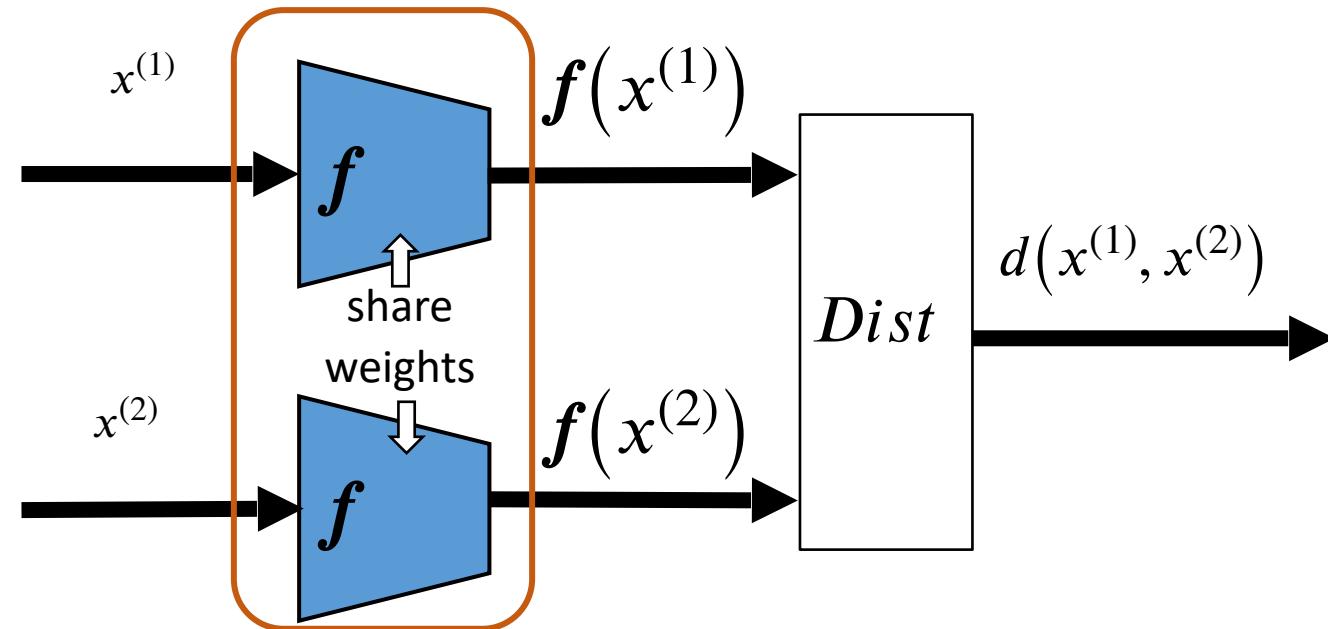
- Two identical networks that share their weights.
- Distance function (e.g., Euclidean)

$$d(x^{(1)}, x^{(2)}) = \|f(x^{(1)}) - f(x^{(2)})\|^2$$



# Training a Siamese network

Update the weights for each branch independently,  
and average them.



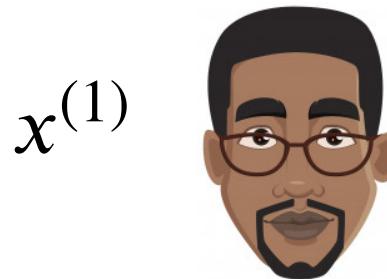
# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks – Contrastive Loss
- Improving Similarity Learning
- Change Detection

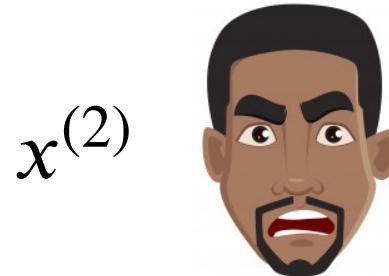
# Contrastive Loss

For a **similar (positive)** pair

- If  $A$  and  $B$  are from the same person,  $d(A, B)$  should be small



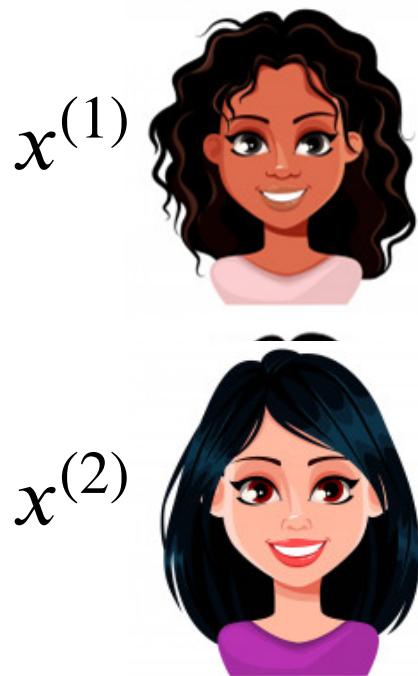
$$\mathcal{L}(x^{(1)}, x^{(2)}) = \|f(x^{(1)}) - f(x^{(2)})\|^2$$



# Contrastive Loss

For a **dissimilar (negative)** pair

- If  $A$  and  $B$  are from different persons,  $d(A, B)$  should be **large**
- and set the loss to zero if  $d(A, B)$  exceeds a margin  $m^2$



$$\begin{aligned} & \left\| f(x^{(1)}) - f(x^{(2)}) \right\|^2 > m^2, \text{ for some } m \\ & m^2 < \left\| f(x^{(1)}) - f(x^{(2)}) \right\|^2 \\ & m^2 - \left\| f(x^{(1)}) - f(x^{(2)}) \right\|^2 < 0 \end{aligned}$$

margin

$$\mathcal{L}(x^{(1)}, x^{(2)}) = \max\left(0, m^2 - \left\| f(x^{(1)}) - f(x^{(2)}) \right\|^2\right)$$

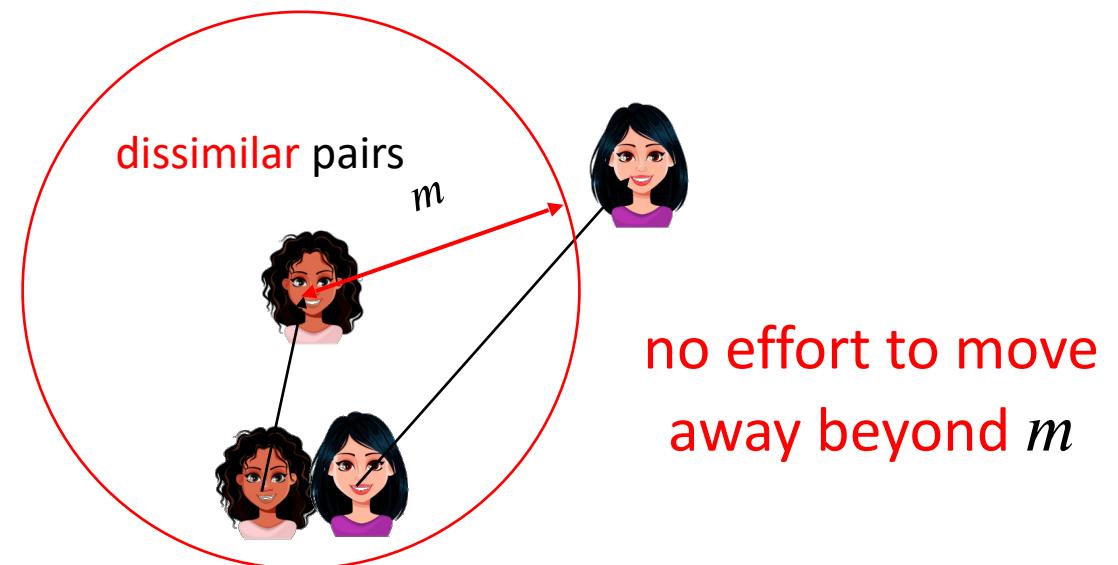
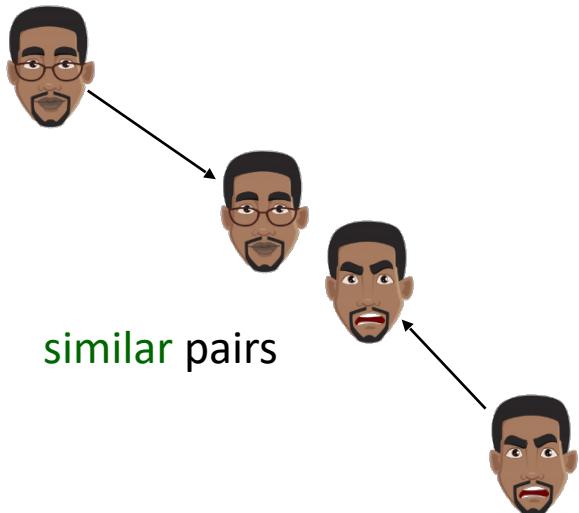
aka Hinge loss

# Contrastive Loss

The general formulation for **any single pair**  $(x^{(1)}, x^{(2)})$  :

$$\mathcal{L} = y \left\| f(x^{(1)}) - f(x^{(2)}) \right\|^2 + (1 - y) \max \left( 0, m^2 - \left\| f(x^{(1)}) - f(x^{(2)}) \right\|^2 \right)$$

↑  
 $y = 0$  for **dissimilar** pairs  
 $y = 1$  for **similar** pairs



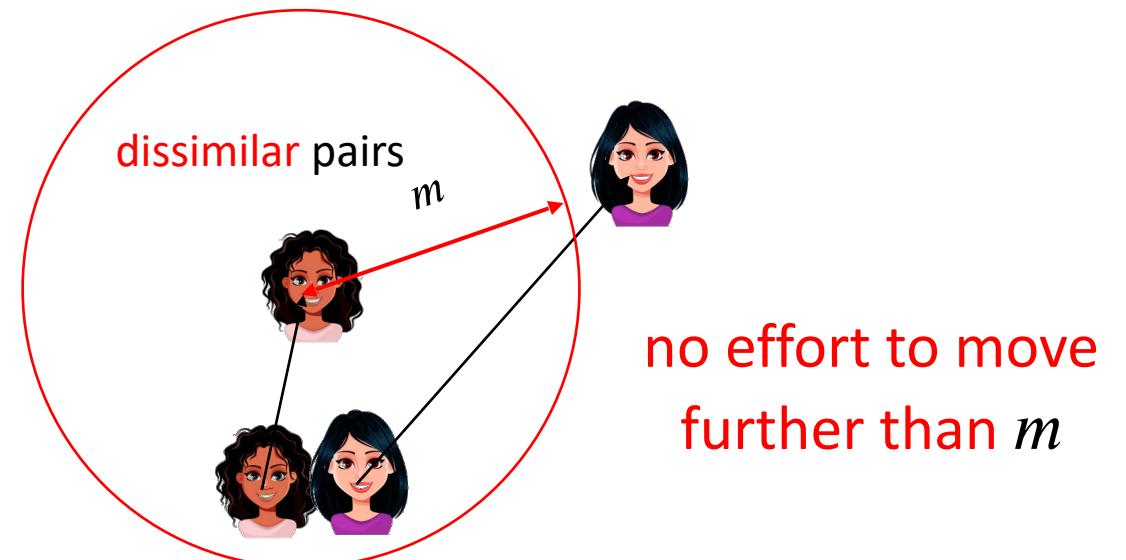
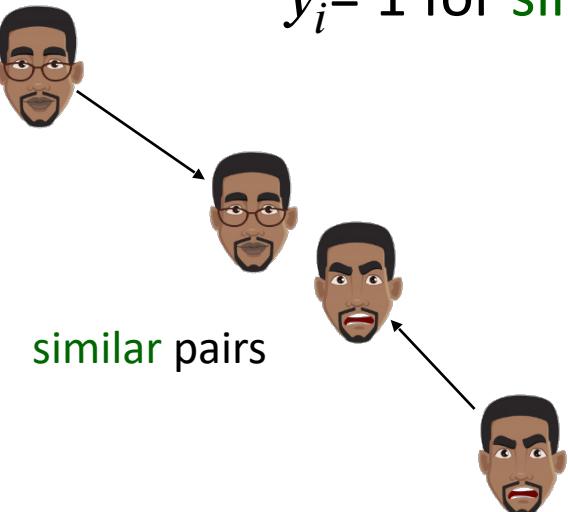
# Contrastive Loss

Formulation considering **all** pairs  $\left\{ (x_i^1, x_i^2) \right\}$  in a batch :

$$\mathcal{L} = \sum_i y_i \left\| f(x_i^{(1)}) - f(x_i^{(2)}) \right\|^2 + (1 - y_i) \max\left( 0, m^2 - \left\| f(x_i^{(1)}) - f(x_i^{(2)}) \right\|^2 \right)$$

↑  
 $y_i = 0$  for **dissimilar** pairs

↑  
 $y_i = 1$  for **similar** pairs



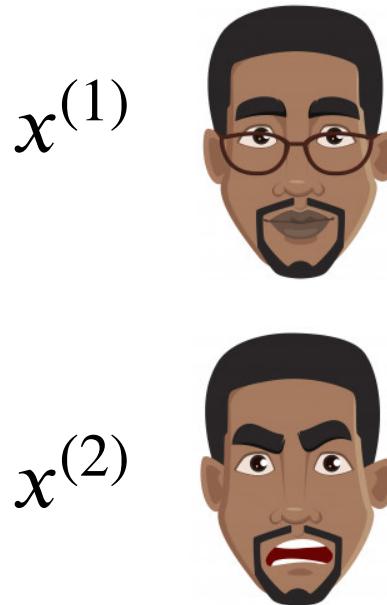
# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks - WDMC
- Improving Similarity Learning
- Change Detection

# Weighted Double Margin Contrastive Loss

For a **similar (positive)** pair

- If  $A$  and  $B$  are from the same person,  $d(A, B)$  should be small
- and set the loss to zero if  $d(A, B)$  is lower than a margin  $m_1^2$

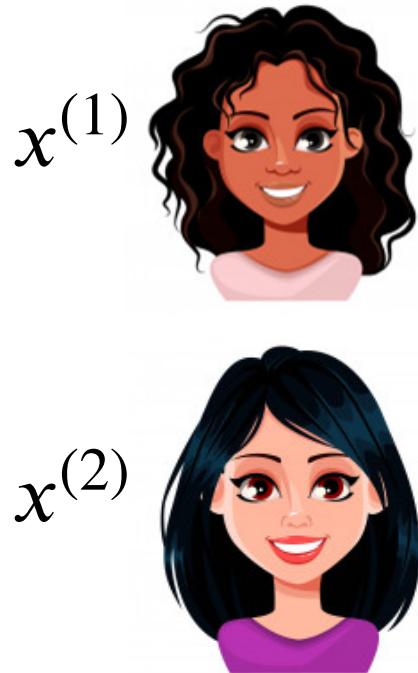


$$\mathcal{L}(x^{(1)}, x^{(2)}) = \max\left(0, \left\|f(x^{(1)}) - f(x^{(2)})\right\|^2 - m_1^2\right)$$

# Weighted Double Margin Contrastive Loss

For a **dissimilar (negative)** pair

- If  $A$  and  $B$  are from different persons,  $d(A, B)$  should be **large**
- and set the loss to zero if  $d(A, B)$  exceeds a margin  $m_2^2$



$$\mathcal{L}(x^{(1)}, x^{(2)}) = \max\left(0, m_2^2 - \|f(x^{(1)}) - f(x^{(2)})\|^2\right)$$

exactly the same  
as for the  
contrastive loss

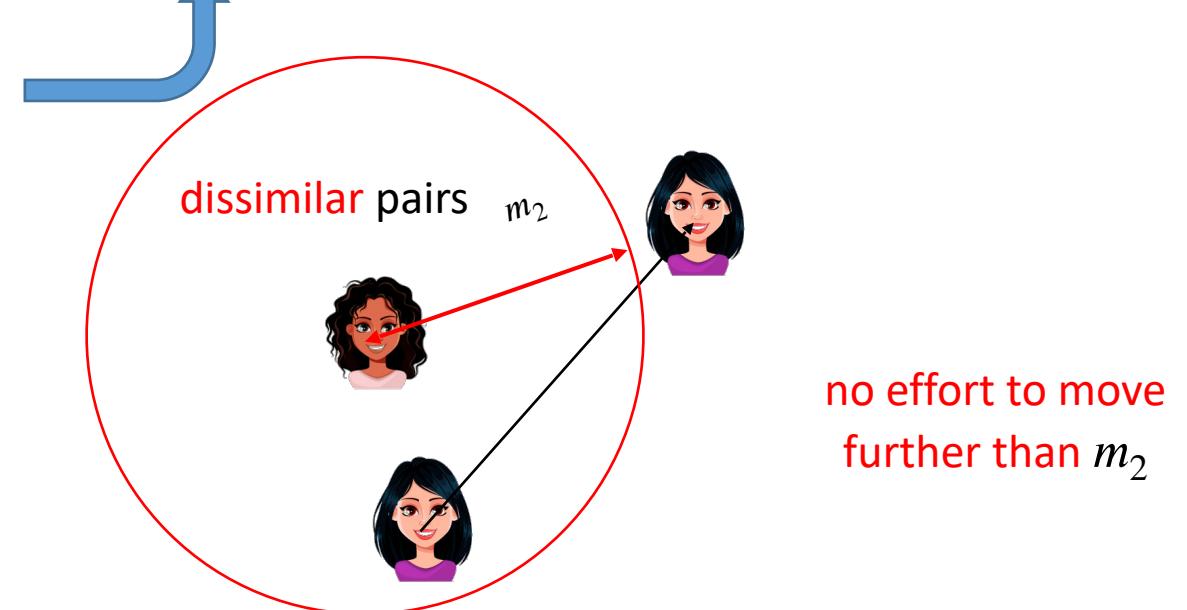
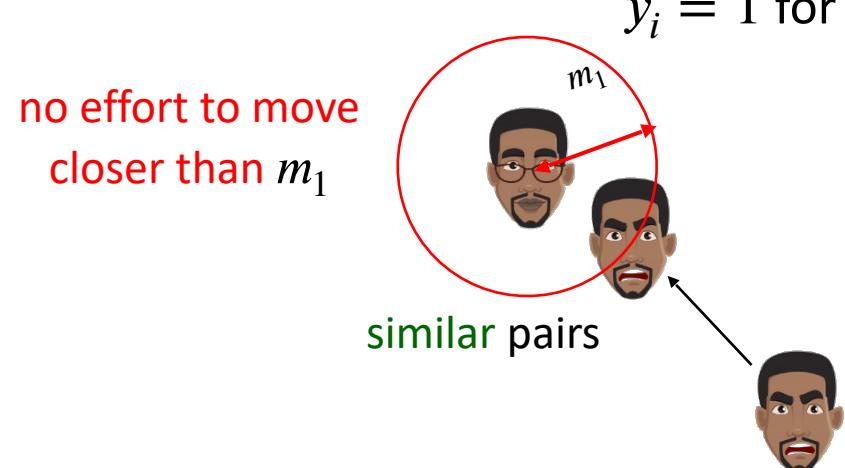
# Weighted Double Margin Contrastive Loss

Formulation considering all pairs  $\left\{ \left( x_i^{(1)}, x_i^{(2)} \right) \right\}$  in a batch :

$$\mathcal{L} = \sum_i w_1 y_i \max\left(0, \left\| f(x_i^{(1)}) - f(x_i^{(2)}) \right\|^2 - m_1^2\right) + w_2 (1 - y_i) \max\left(0, m_2^2 - \left\| f(x_i^{(1)}) - f(x_i^{(2)}) \right\|^2\right)$$

$w_1, w_2$  are the weights for  
similar, dissimilar pairs

$y_i = 0$  for dissimilar pairs  
 $y_i = 1$  for similar pairs



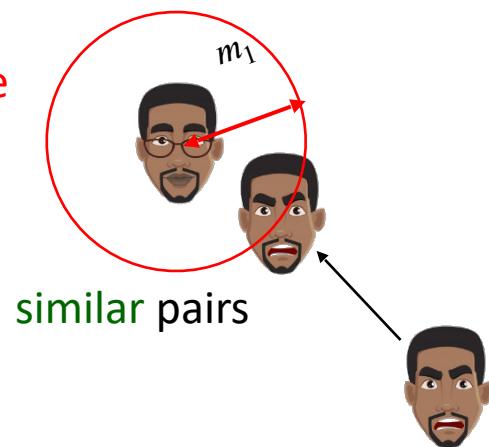
# Weighted Double Margin Contrastive loss

Formulation considering **all** pairs  $\left\{ \left( x_i^{(1)}, x_i^{(2)} \right) \right\}$  in a batch :

$$\mathcal{L} = \sum_i y_i w_1 \max \left( 0, \left\| f(x_i^{(1)}) - f(x_i^{(2)}) \right\|^2 - m_1^2 \right) + (1 - y_i) w_2 \max \left( 0, m_2^2 - \left\| f(x_i^{(1)}) - f(x_i^{(2)}) \right\|^2 \right)$$

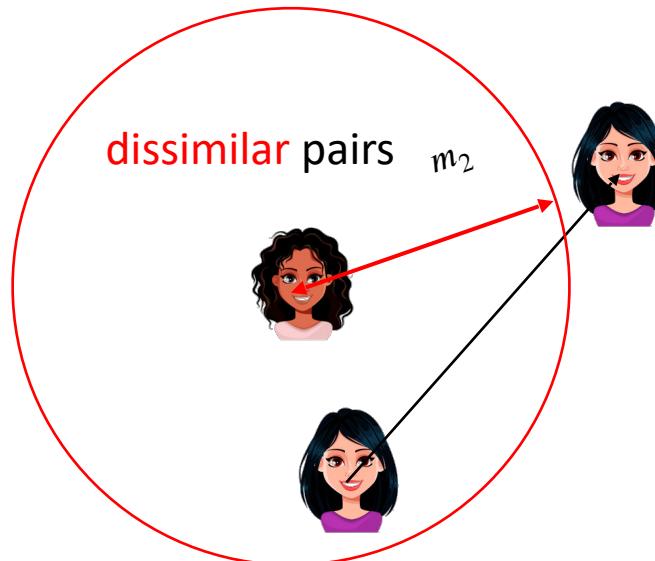
$y_i = 0$  for **dissimilar** pairs  
 $y_i = 1$  for **similar** pairs

no effort to move closer than  $m_1$



dissimilar pairs

no effort to move further than  $m_2$



# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks – Triplet Loss
- Improving Similarity Learning
- Change Detection

# Triplet loss

Training is conducted on triplets instead of pairs:



Anchor ( $A$ )  
any arbitrary sample



Positive ( $P$ )  
same class as the anchor



Negative ( $N$ )  
different class from the anchor

We want  $\|f(A) - f(P)\|^2 < \|f(A) - f(N)\|^2$

# Triplet loss

For a single triplet:

$$\|f(A) - f(\textcolor{green}{P})\|^2 < \|f(A) - f(\textcolor{red}{N})\|^2$$

$$\|f(A) - f(\textcolor{green}{P})\|^2 - \|f(A) - f(\textcolor{red}{N})\|^2 < 0$$

Making it more restrictive

$$\|f(A) - f(\textcolor{green}{P})\|^2 - \|f(A) - f(\textcolor{red}{N})\|^2 < -m$$

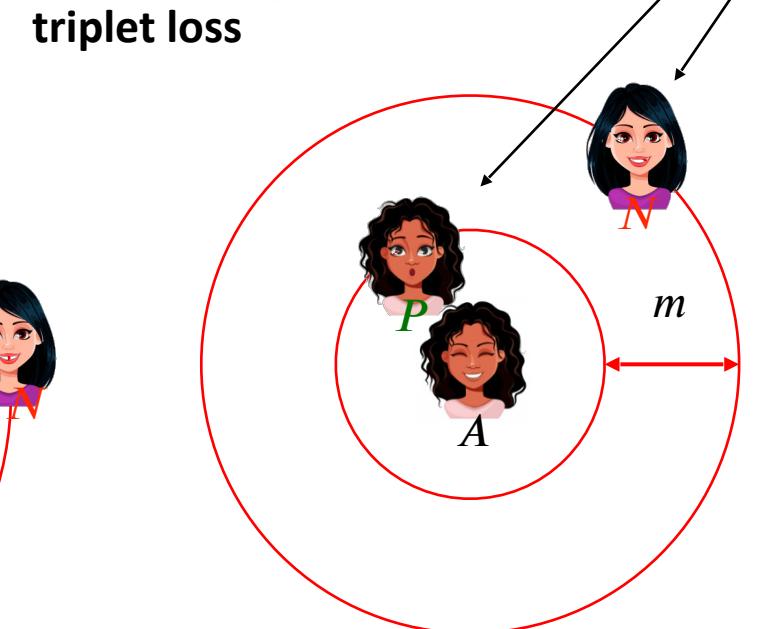
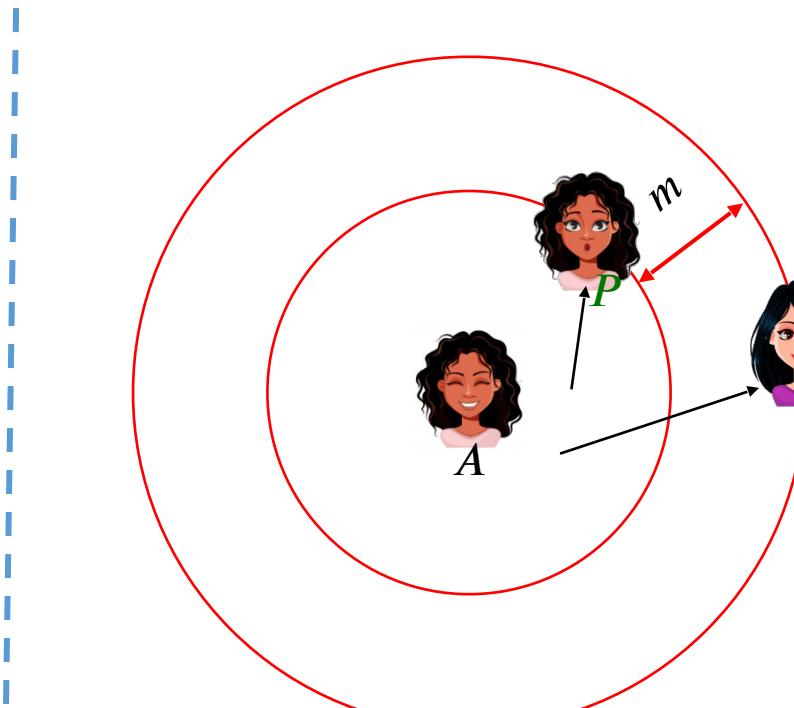
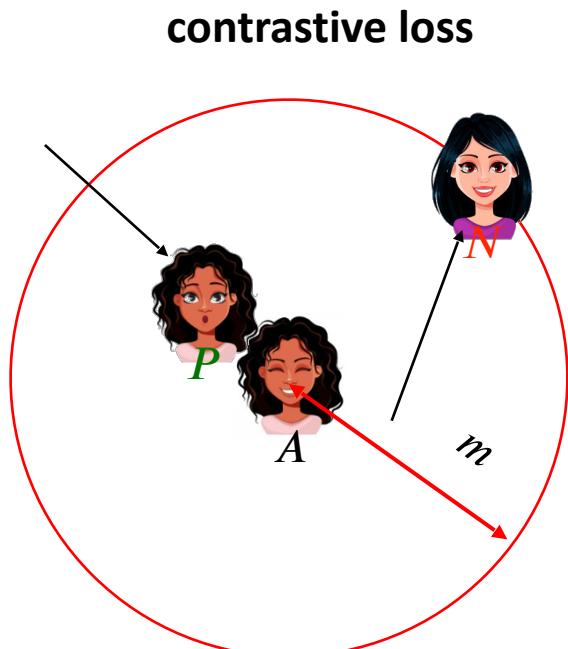
margin

for some  $m > 0$

$$\mathcal{L}(A, \textcolor{green}{P}, \textcolor{red}{N}) = \max\left(0, \|f(A) - f(\textcolor{green}{P})\|^2 - \|f(A) - f(\textcolor{red}{N})\|^2 + m\right)$$

# Contrastive vs. Triplet Loss

For a single triplet:



what about  
WDMC?

# Triplet loss

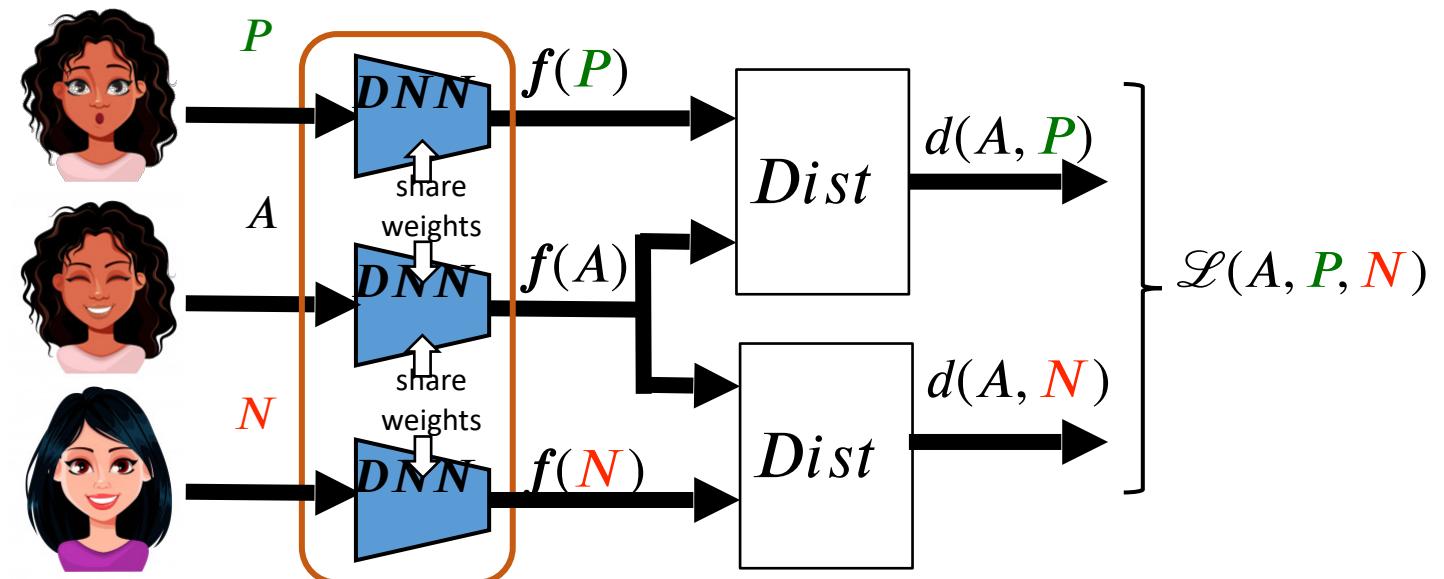
Formulation considering **all triplets**  $\{ (A_i, \textcolor{red}{P}_i, \textcolor{green}{N}_i) \}$  in a batch:

$$\mathcal{L} = \sum_i \max \left( \left( 0, \|f(A_i) - f(\textcolor{red}{P}_i)\|^2 - \|f(A_i) - f(\textcolor{green}{N}_i)\|^2 + m \right) \right)$$

# Triplet architecture

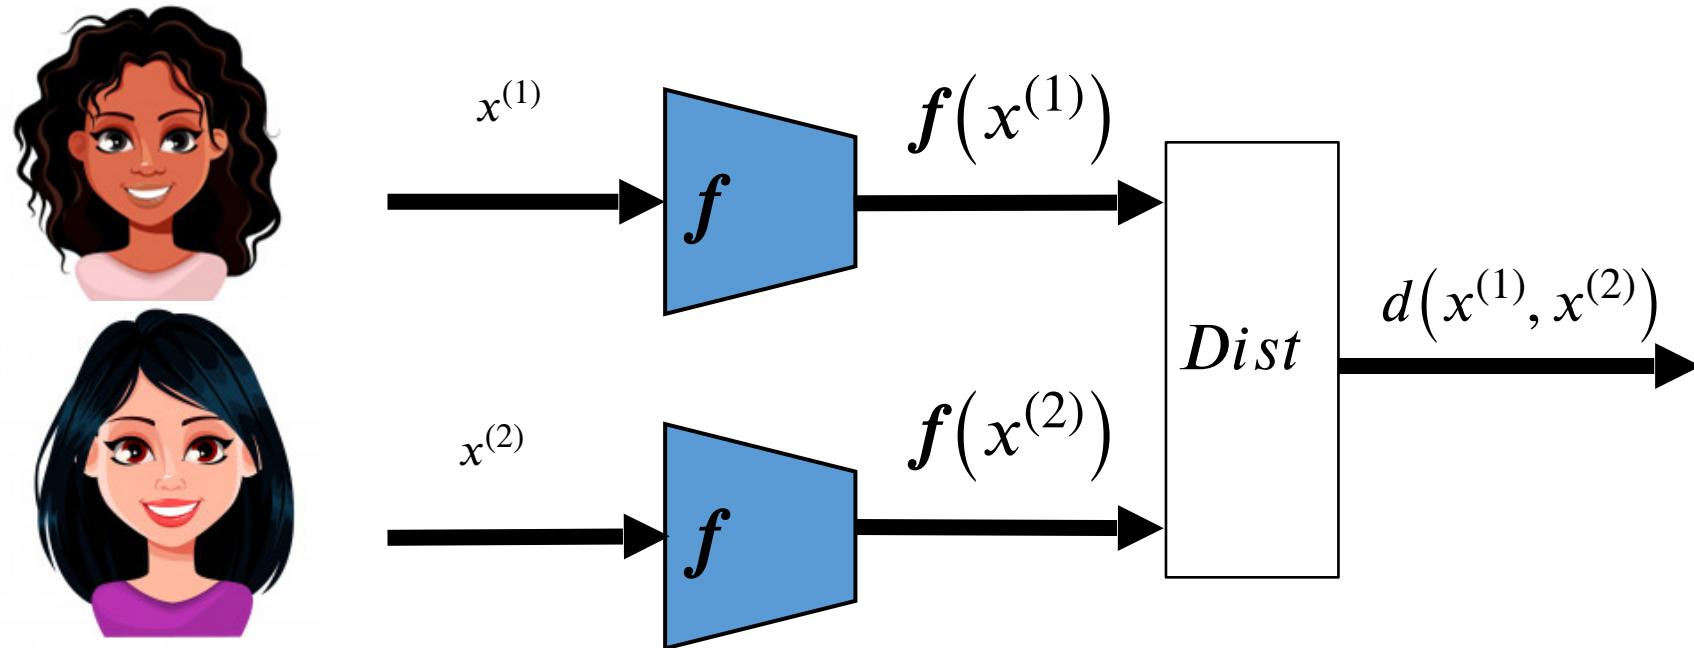
## Training

- 3-branch network
- Positive & negative branches
- Always share weights

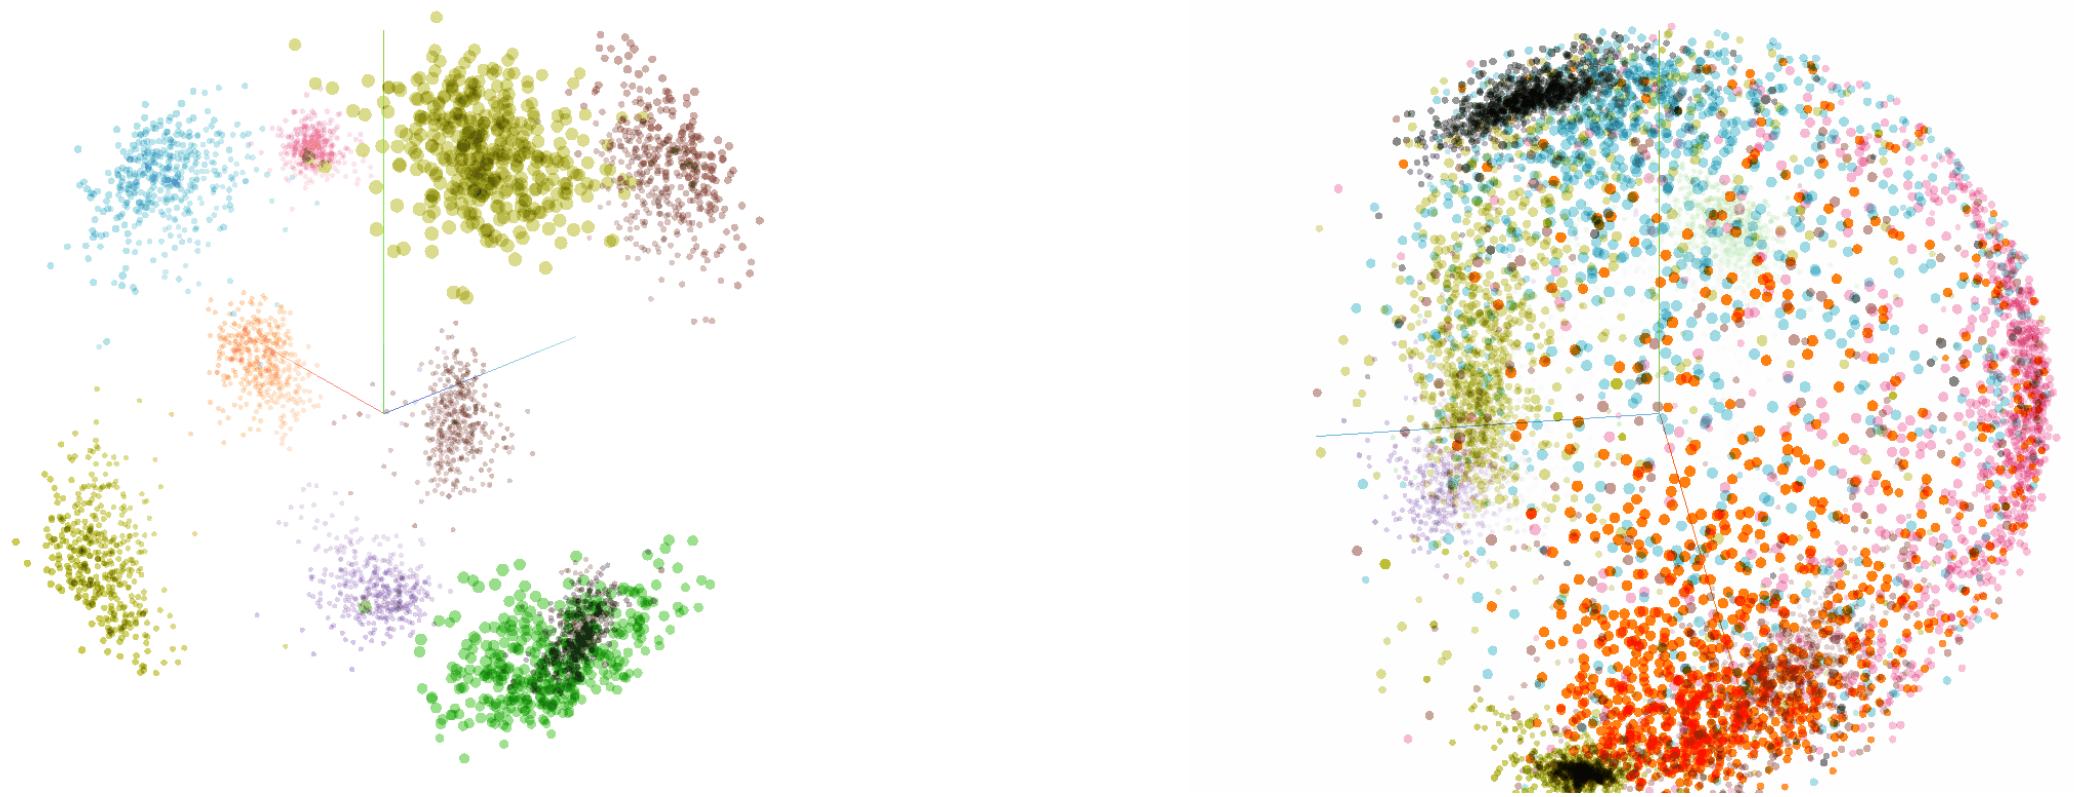


# Triplet architecture

Test: same as for the contrastive case.



# View of the Embedding Space



Embedding Space for train data (left) and test data (right) using PCA in Tensorboard

# Test time

## Just nearest neighbor search



Query

Often in security applications, the subjects in  $n$  highest rank positions are analyzed visually by a human expert for a final decision.

Ranking

1<sup>st</sup>



2<sup>nd</sup>

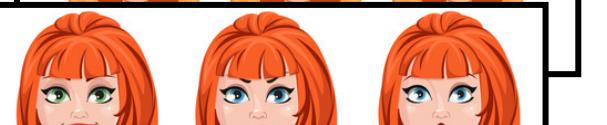
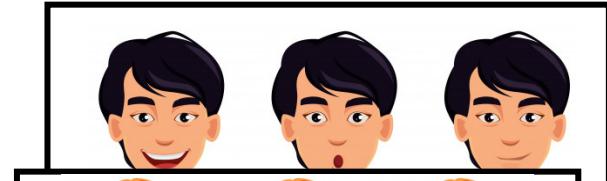


3<sup>rd</sup>



⋮

Gallery



identity

# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks
- Improving Similarity Learning
- Change Detection

# Improving Similarity Learning

- Loss:
  - Contrastive vs. triplet loss vs. WDMC
- Sampling
  - Sample the space wisely diversity of classes & hard cases
- Ensembles
  - Multiple networks, each trained with a subset of triplets
- Use the Classification loss for similarity learning

# References for Losses

- Wang et al., 2017 Deep metric learning with angular loss.
- Yu et al., 2018, Correcting the triplet selection bias for triplet loss

# Improving Similarity Learning

- Loss:
  - Contrastive vs. triplet loss vs. WDMC
- Sampling
  - Sample the space wisely exploiting diversity of classes & hard cases
- Ensembles
  - Multiple networks, each trained with a subset of triplets
- Use the Classification loss for similarity learning

# Problems with the Triplet Loss

1. The number of possible triplets is  $O(n^3)$ . Training with all of them would require a very long time.
2. Recall that

- the triplet loss  $\rightarrow \mathcal{L}(A, P, N) = \max(0, d(A, P)^2 - d(A, N)^2 + m)$

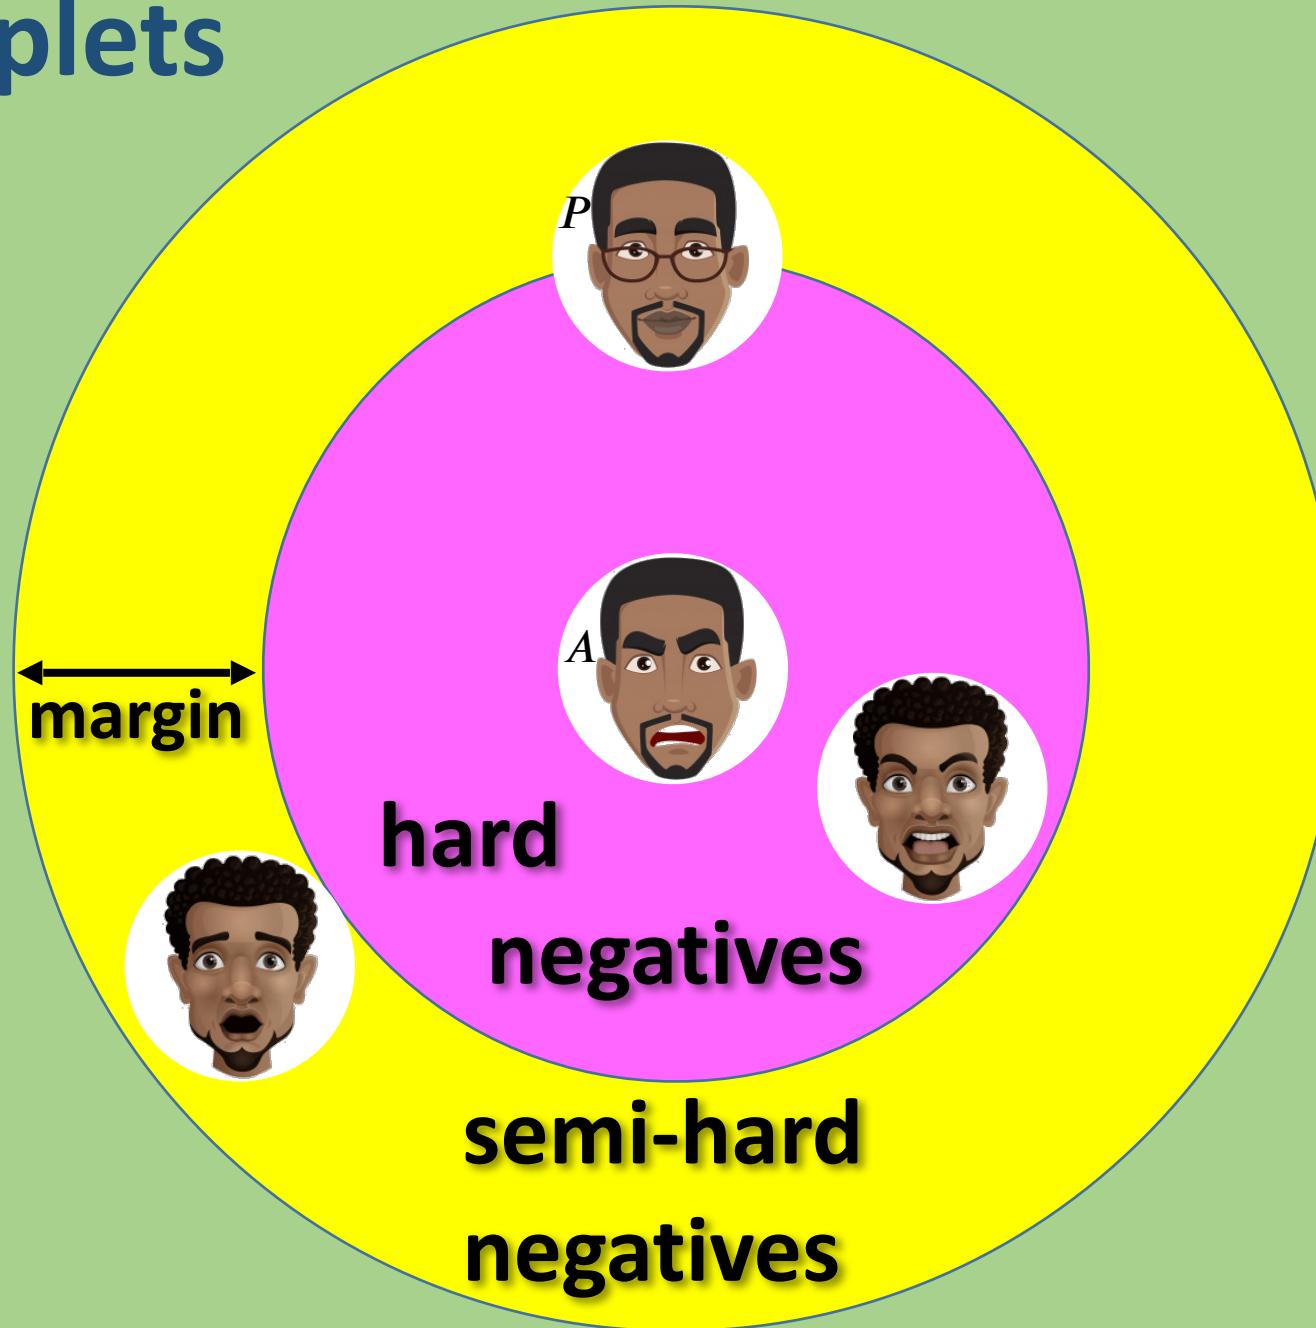
- the gradient of  $\max$  function  $\rightarrow \frac{d \max(0, x)}{dx} = \begin{cases} 0 & \text{if } x \leq 0 \\ 1 & \text{if } x > 0 \end{cases}$

Therefore, triplets for which

$$\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + m \leq 0$$

will not contribute to parameter learning.

# Choosing triplets



**easy  
negatives**

**no contribution  
to learning**

# Hard triplet mining

General procedure:

```
choose initial_triplets { $A_i, P_i, N_i$ },
```

```
while stop_condition = false do
```

- Train for some epochs

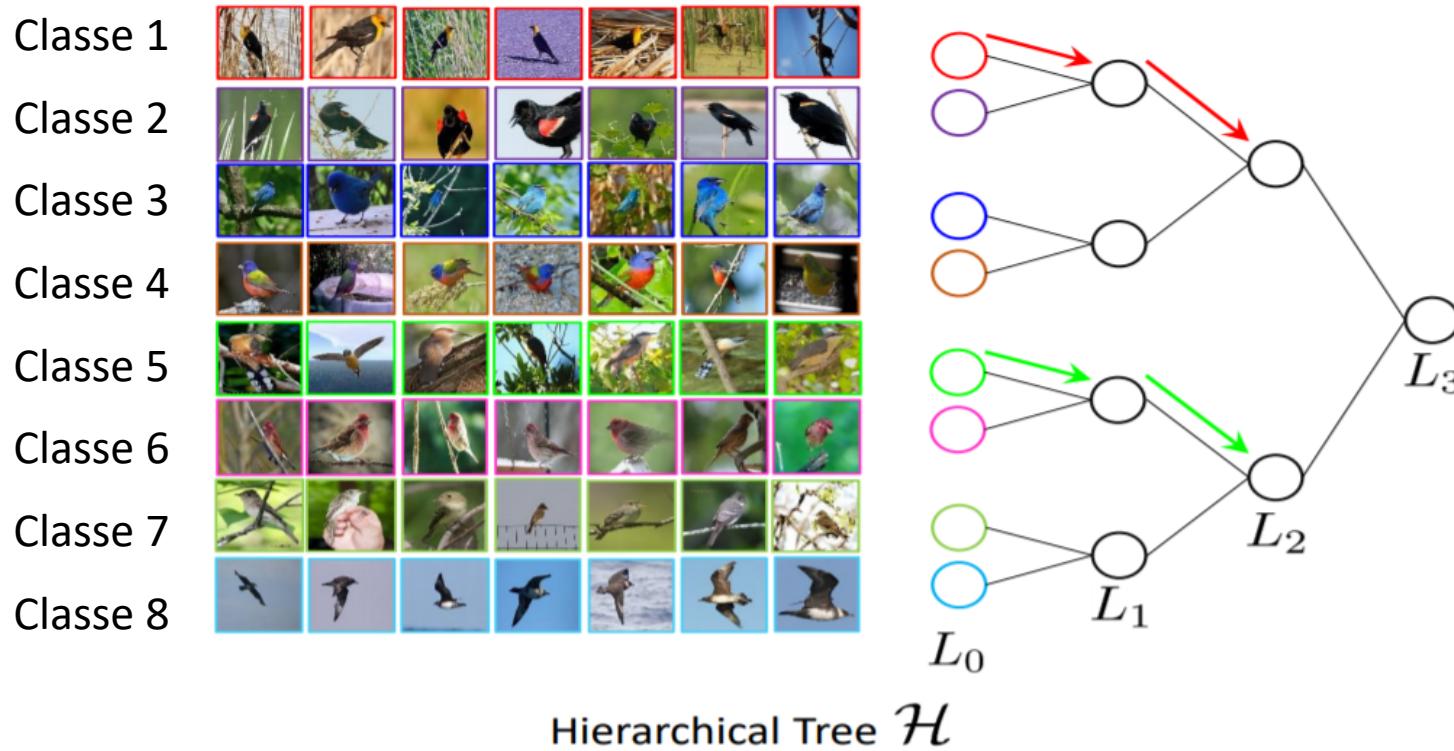
- Choose hard triplets { $A_i, P_i, N_i$ }

```
end
```

Hard triplets can be selected offline (manually) or online (automatically)

# Sampling: Hierarchical Triplet Loss

1. Build a hierarchical tree where the leaves represent classes.
2. Recursively merge them until the root node



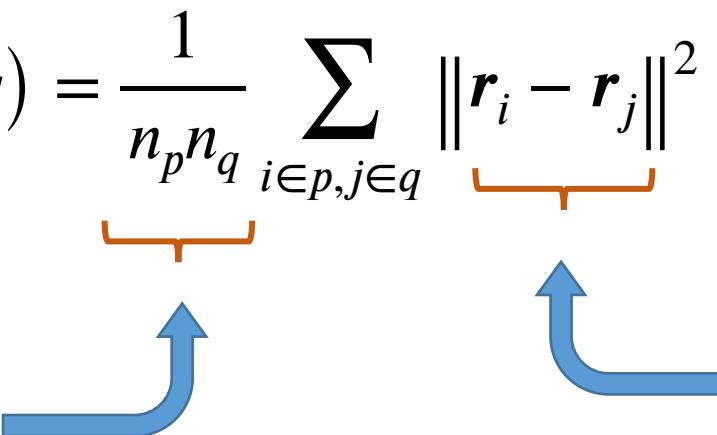
# HTL: Building the Tree

- To create the tree, first define the **between class distances**  $d(p, q)$ :

$$d(p, q) = \frac{1}{n_p n_q} \sum_{i \in p, j \in q} \| \mathbf{r}_i - \mathbf{r}_j \|^2$$

number of samples in classes  $p$  and  $q$

deep features of samples  $i$  and  $j$



- If the deep features  $\mathbf{r}_i$  are normalized into unit length,  $d(p, q)$  varies from 0 to 4.

# HTL: Building the Tree

- Calculate the **average within classes distance**  $d_0$ :

$$d_0 = \frac{1}{C} \sum_{c=1}^C \left( \frac{1}{n_c^2 - n_c} \sum_{i,j \in c} \|r_i - r_j\|^2 \right)$$

number of classes

twice the number of sample pairs in class  $c$

number of sample in class  $c$

distance between samples  $i$  and  $j$  in the embedding space

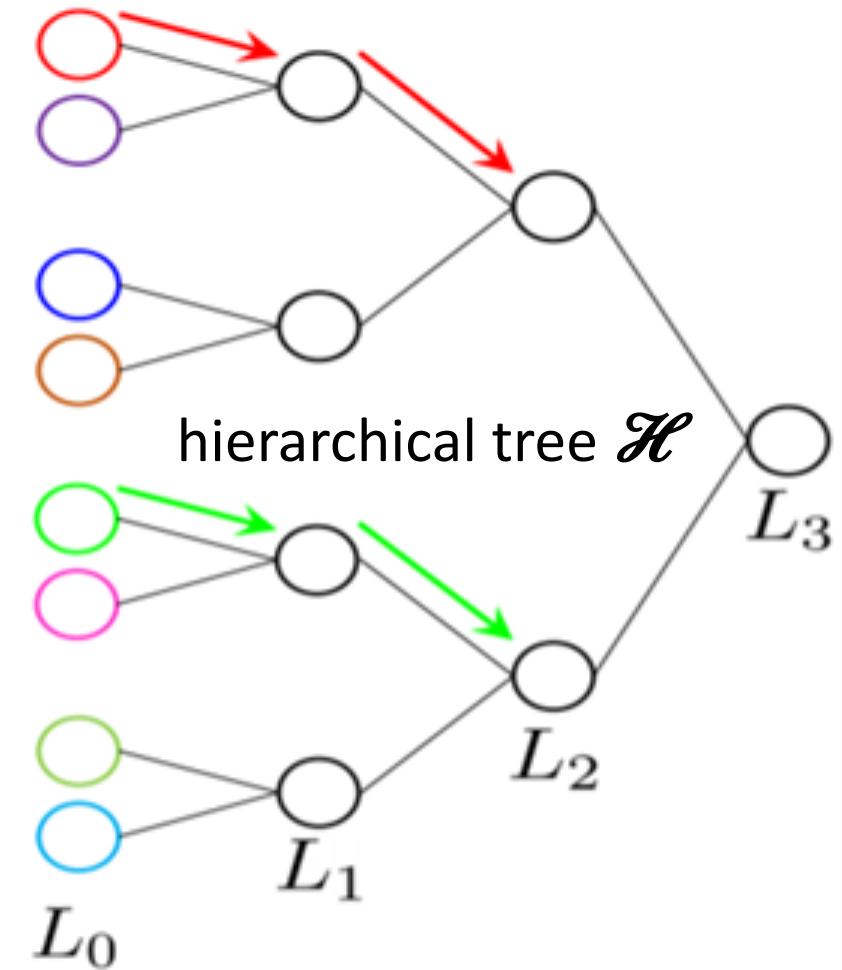
# HTL: Building the Tree

- Two classes with a distance less than

$$d_l = \frac{l(4 - d_0)}{L} + d_0$$

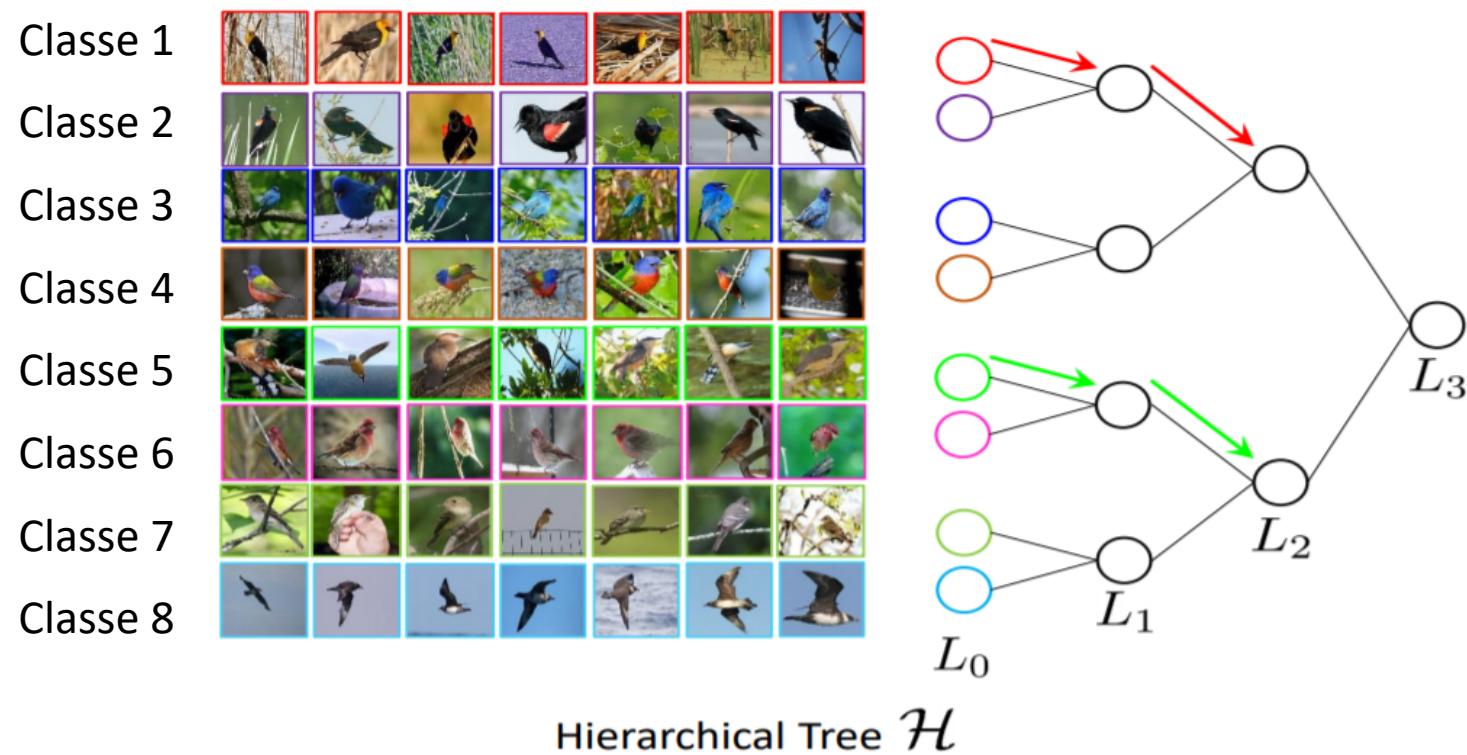
are merged into a single node at the  $l$ -th level.

- Nodes are merged from 0-th level to a single node at to  $L$ -th level forming the hierarchical tree  $\mathcal{H}$ .



# HTL: Anchor-Neighbor Sampling

- Randomly select  $l'$  nodes at 0<sup>th</sup> level
    - Preserve class diversity in the mini-batch



# HTL: Anchor-Neighbor Sampling

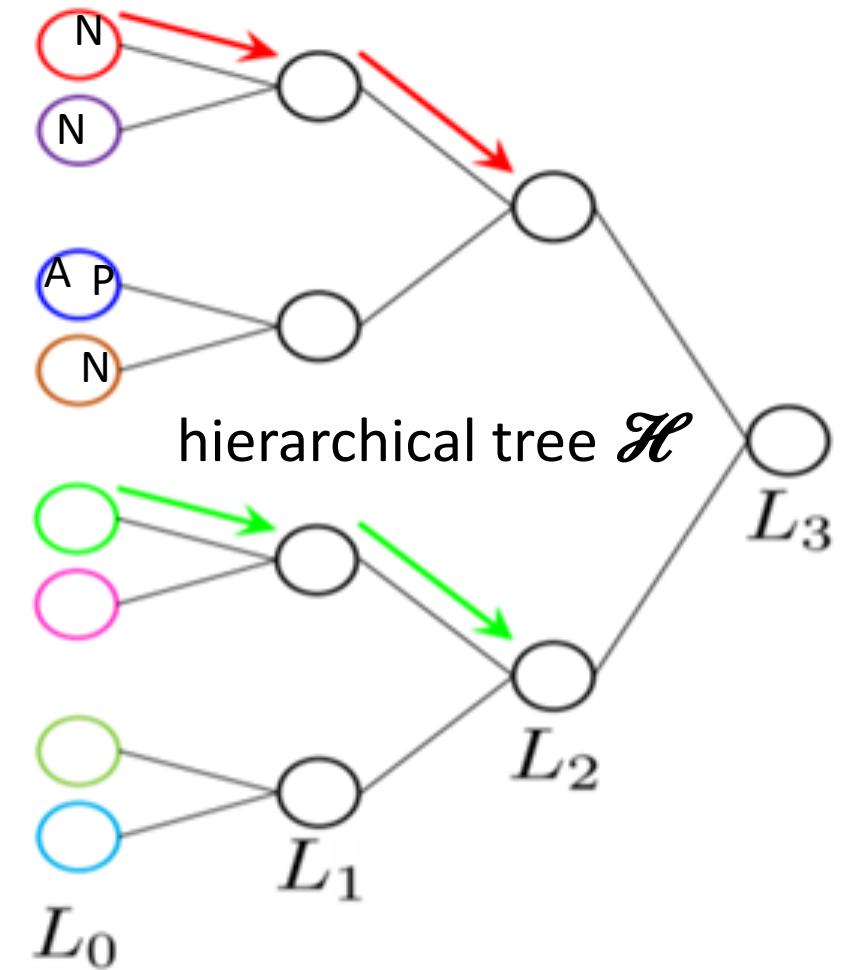
- Randomly select  $l'$  nodes at 0<sup>th</sup> level, and select one **anchor** sample from each node
- For each anchor select the furthest **positive** sample.
- For each of the  $l'$  anchors, pick the  $k$  nearest nodes/classes at 0<sup>th</sup> level based on the class distance measured in the feature space.
- Collect  $t$  samples from these nodes/classes randomly to be **negative** samples.

We will end up with a mini batch containing  $l' k t$  triplets.

# HTL: Anchor-Neighbor Sampling

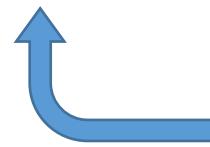
- Randomly select  $l'$  nodes at 0<sup>th</sup> level, and select one **anchor** sample from each node
- For each anchor select the furthest **positive** sample.
- For each of the  $l'$  anchors, pick the  $k$  nearest nodes/classes at 0<sup>th</sup> level based on the class distance measured in the feature space.
- Collect  $t$  samples from these nodes/classes randomly to be **negative** samples.

We will end up with a mini batch containing  $l' k t$  triplets.



# HTL: Loss Formulation

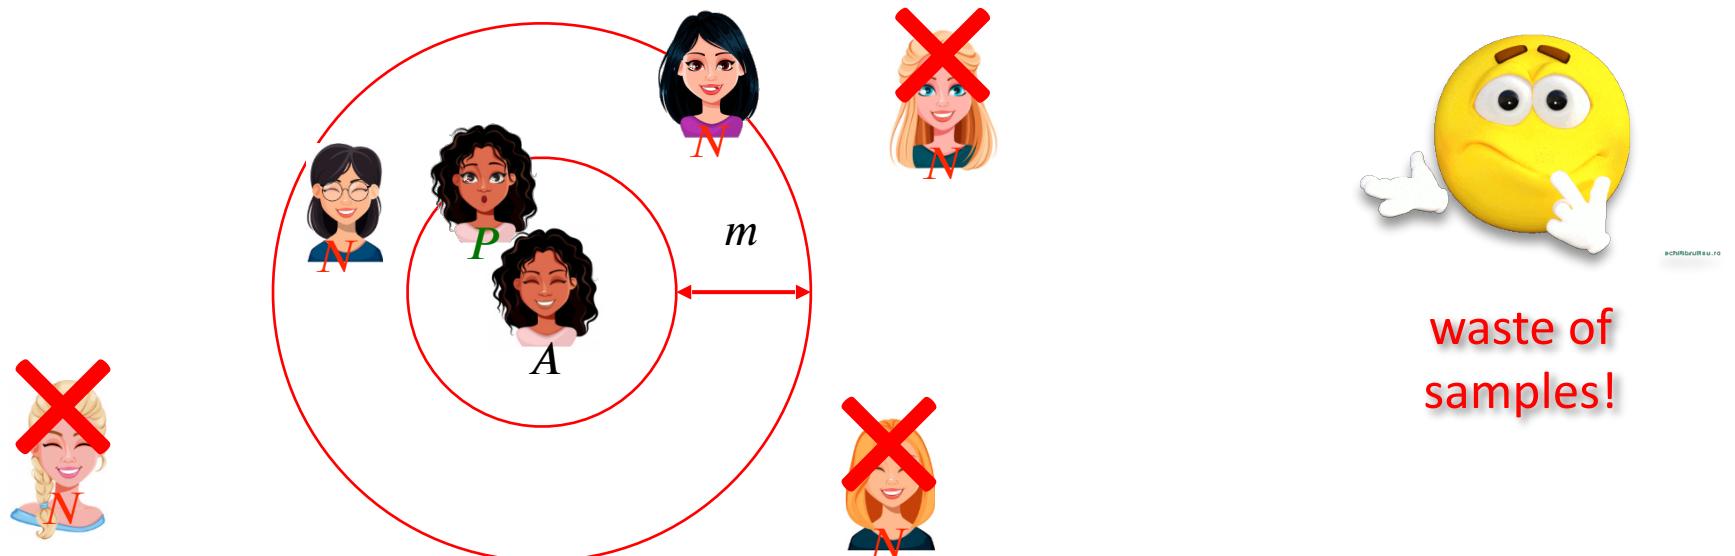
$$\mathcal{L} = \sum_i \max\left(0, \left\|f(\mathbf{x}_A^i) - f(\mathbf{x}_P^i)\right\|^2 - \left\|(\mathbf{x}_A^i) - f(\mathbf{x}_N^i)\right\|^2 + m\right)$$



all the triplets in the batch

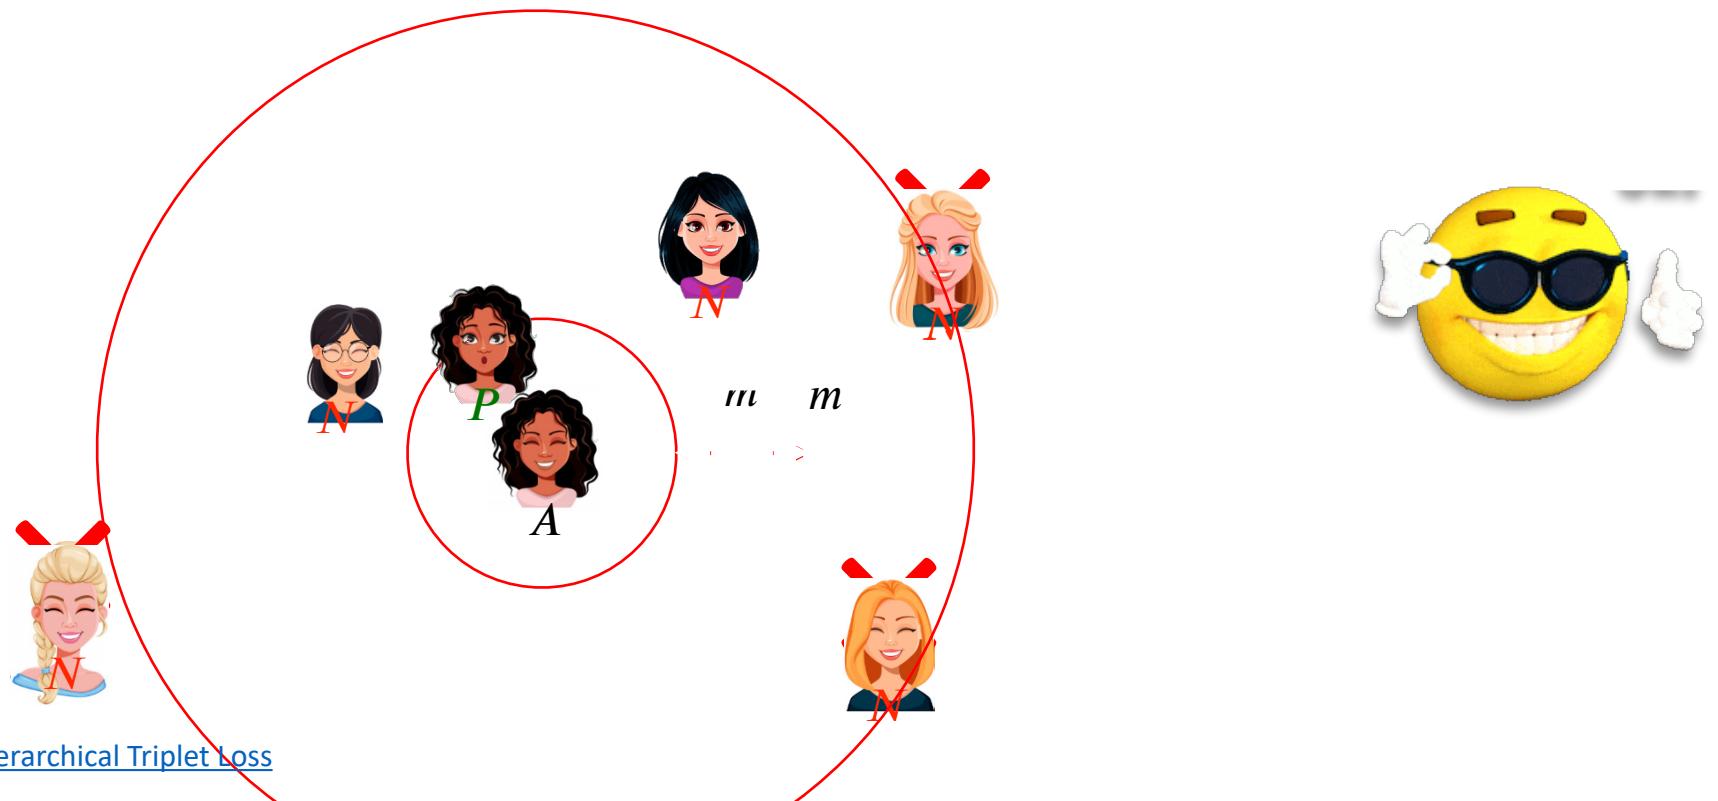
# Dynamic Margin

- Hard negative mining rejects negatives “too far” from the anchor, whereby “too far” is defined by the margin  $m$ .



# Dynamic Margin

- “*... to learn from more meaningful hard samples with the guide of global context,*” we keep the distant negatives and enlarge the margin.



# Dynamic Margin

... the loss takes the form

$$\mathcal{L} = \sum_i \max\left(0, \left\| f(\mathbf{x}_A^i) - f(\mathbf{x}_P^i) \right\|^2 - \left\| (\mathbf{x}_A^i) - f(\mathbf{x}_N^i) \right\|^2 + m_i\right)$$

The margin  $m_i$  depends on the distances between the anchor and the negative classes computed in the hierarchical tree (see ref. for details).

The distances change as training progresses. Therefore, the tree must be recalculated after some epochs.

# HTL Performance

## Example:

Results on the *In-Shop Clothes Retrieval* [dataset](#)

The accuracy metric is the ***Recall@K***.

*“If one of the K retrieved images has the same label as the query, the recall will increase by 1.”*



# HTL Performance

## Accuracy

Comparisons on the In-Shop Clothes Retrieval dataset

R@	1	10	20	30	40	50
FashionNet+Joints[15]	41.0	64.0	68.0	71.0	73.0	73.5
FashionNet+Poselets[15]	42.0	65.0	70.0	72.0	72.0	75.0
FashionNet[15]	53.0	73.0	76.0	77.0	79.0	80.0
HDC[38]	62.1	84.9	89.0	91.2	92.3	93.1
BIER[18]	76.9	92.8	95.2	96.2	96.7	97.1
Triplet loss →	62.3	85.1	89.0	91.1	92.4	93.4
Anchor-Neighbor →	75.3	91.8	94.3	96.2	96.7	97.5
With dynamic margin →	<b>80.9</b>	<b>94.3</b>	<b>95.8</b>	<b>97.2</b>	<b>97.4</b>	<b>97.8</b>

# References for Sampling

- Manmatha et al., 2017, Sampling matters for deep metric learning.
- Xu et al., 2019, Deep asymmetric learning via rich relationship mining.
- Duan et al., 2019, Deep embedding learning with discriminative sampling policy
- Wang et al., 2019, Multi-similarity loss with general pair weighting for deep metric learning.
- Wang et al., 2019, Ranked list loss for deep metric learning.

# Improving Similarity Learning

- Loss:
  - Contrastive vs. triplet loss vs. WDMC
- Sampling
  - Sample the space wisely diversity of classes & hard cases
- Ensembles
  - Multiple networks, each trained with a subset of triplets
- Use the Classification loss for similarity learning

# Ensembles: Divide and Conquer

## Rational:

- The goal is learning an embedding space where semantically similar objects are close and dissimilar objects are far apart.
- We have been after a single distance metric that encodes many different notions of similarity such as color, shape, pose or semantic meaning. This is too much for a single NN.

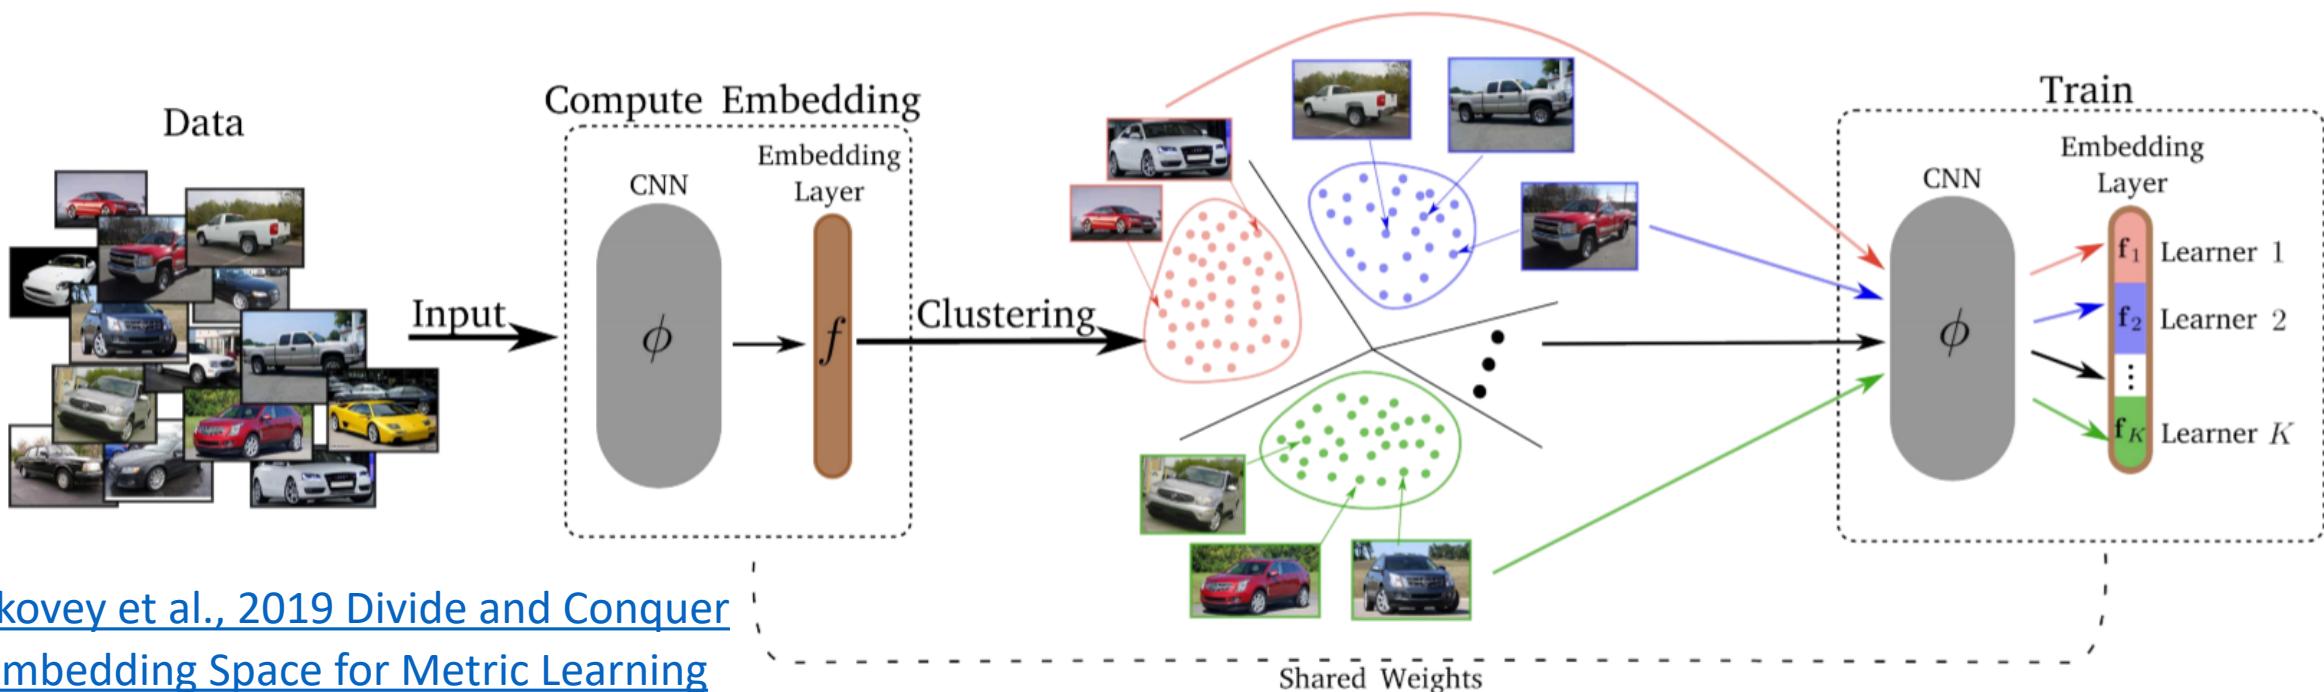
## Intuition:

- Split the problem into simpler sub-problems, train the learners separately, and merge them after they have converged.

# Divide & Conquer

The mapping into the  $d$ -dimensional embedding space is a non linear function  $\phi$  followed by a linear layer  $f$ .

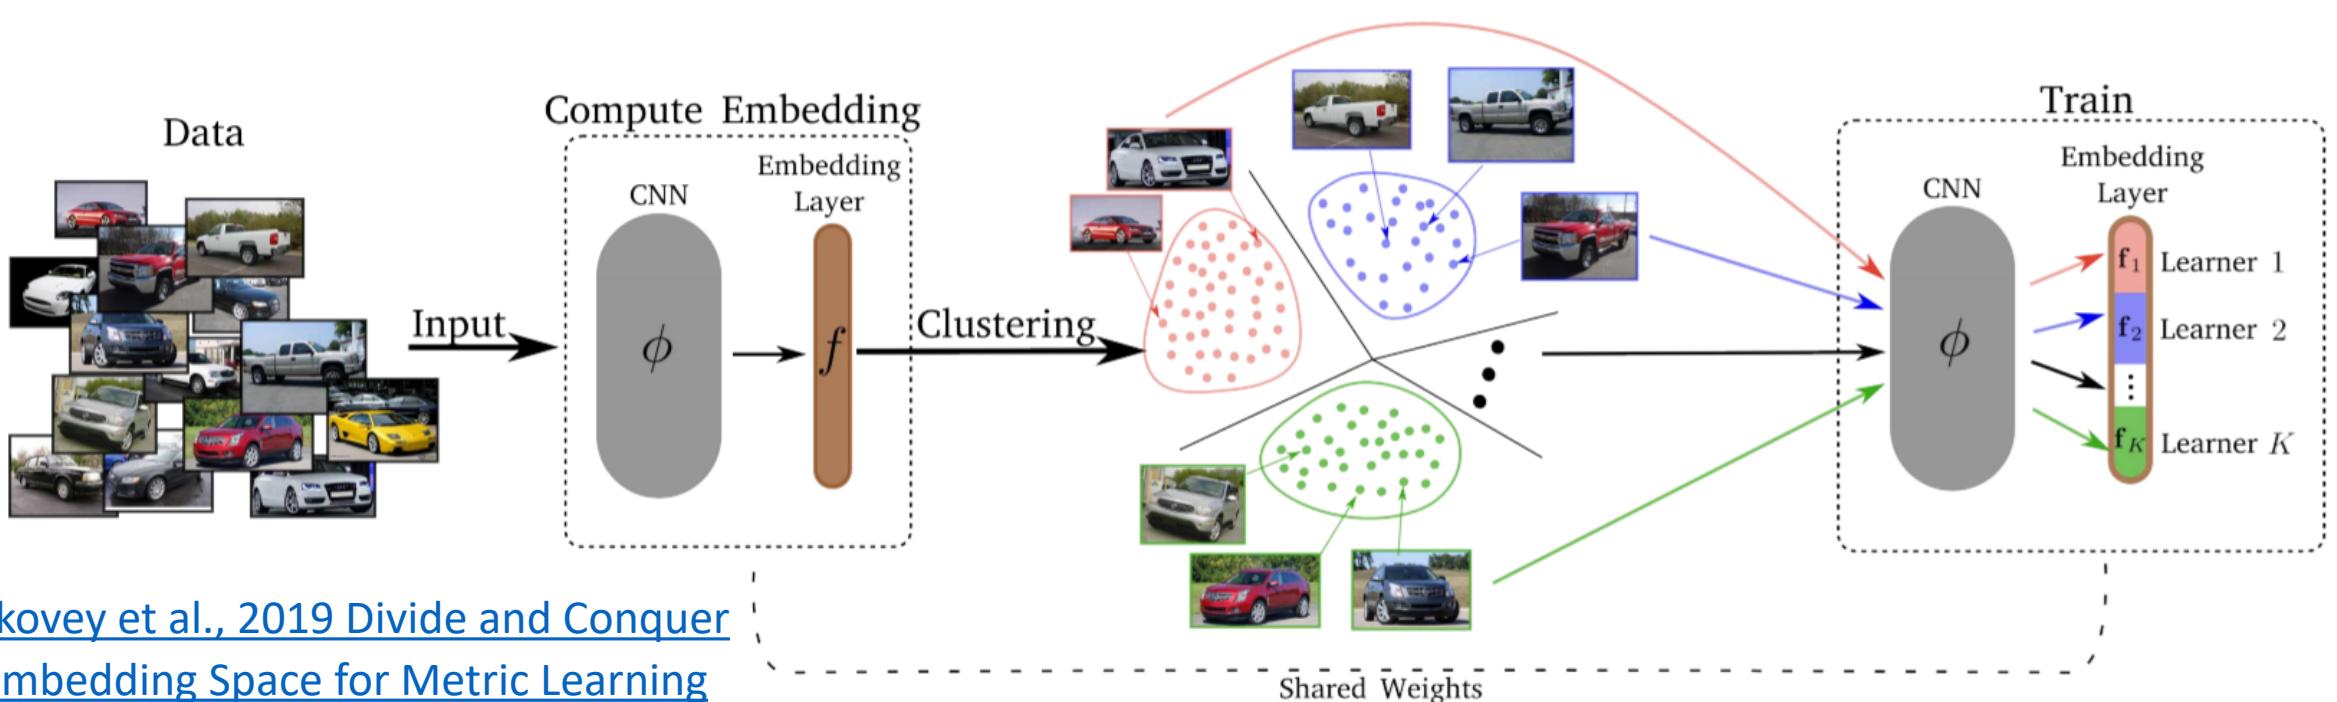
So, representation of an input  $x$  in the embedding space will be the  $d$ -dimensional vector  $f(\phi(x))$



# Divide & Conquer

## DIVIDE

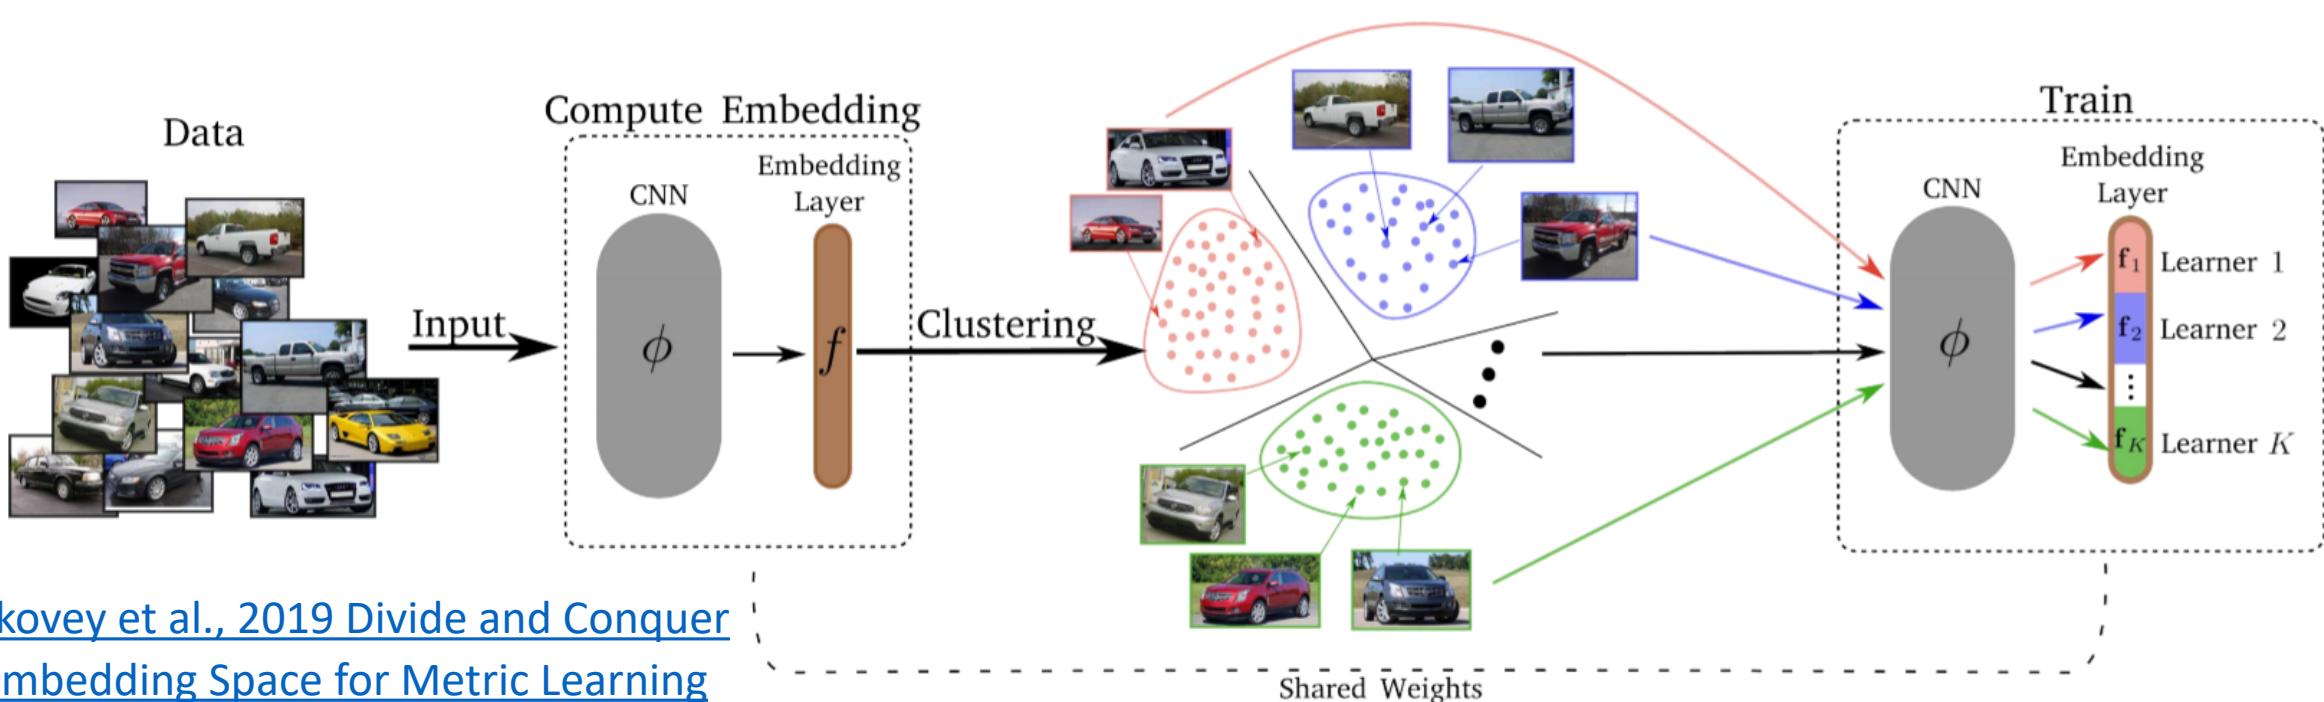
- 1) Cluster the samples in the embedding space into  $K$  cluster using K-means



# Divide & Conquer

## DIVIDE

- 2) Define  $K$  learners consisting of the same non linear function  $\phi$  followed by a linear layer  $f_k$ , that maps the input into a  $d/K$  dimensional space

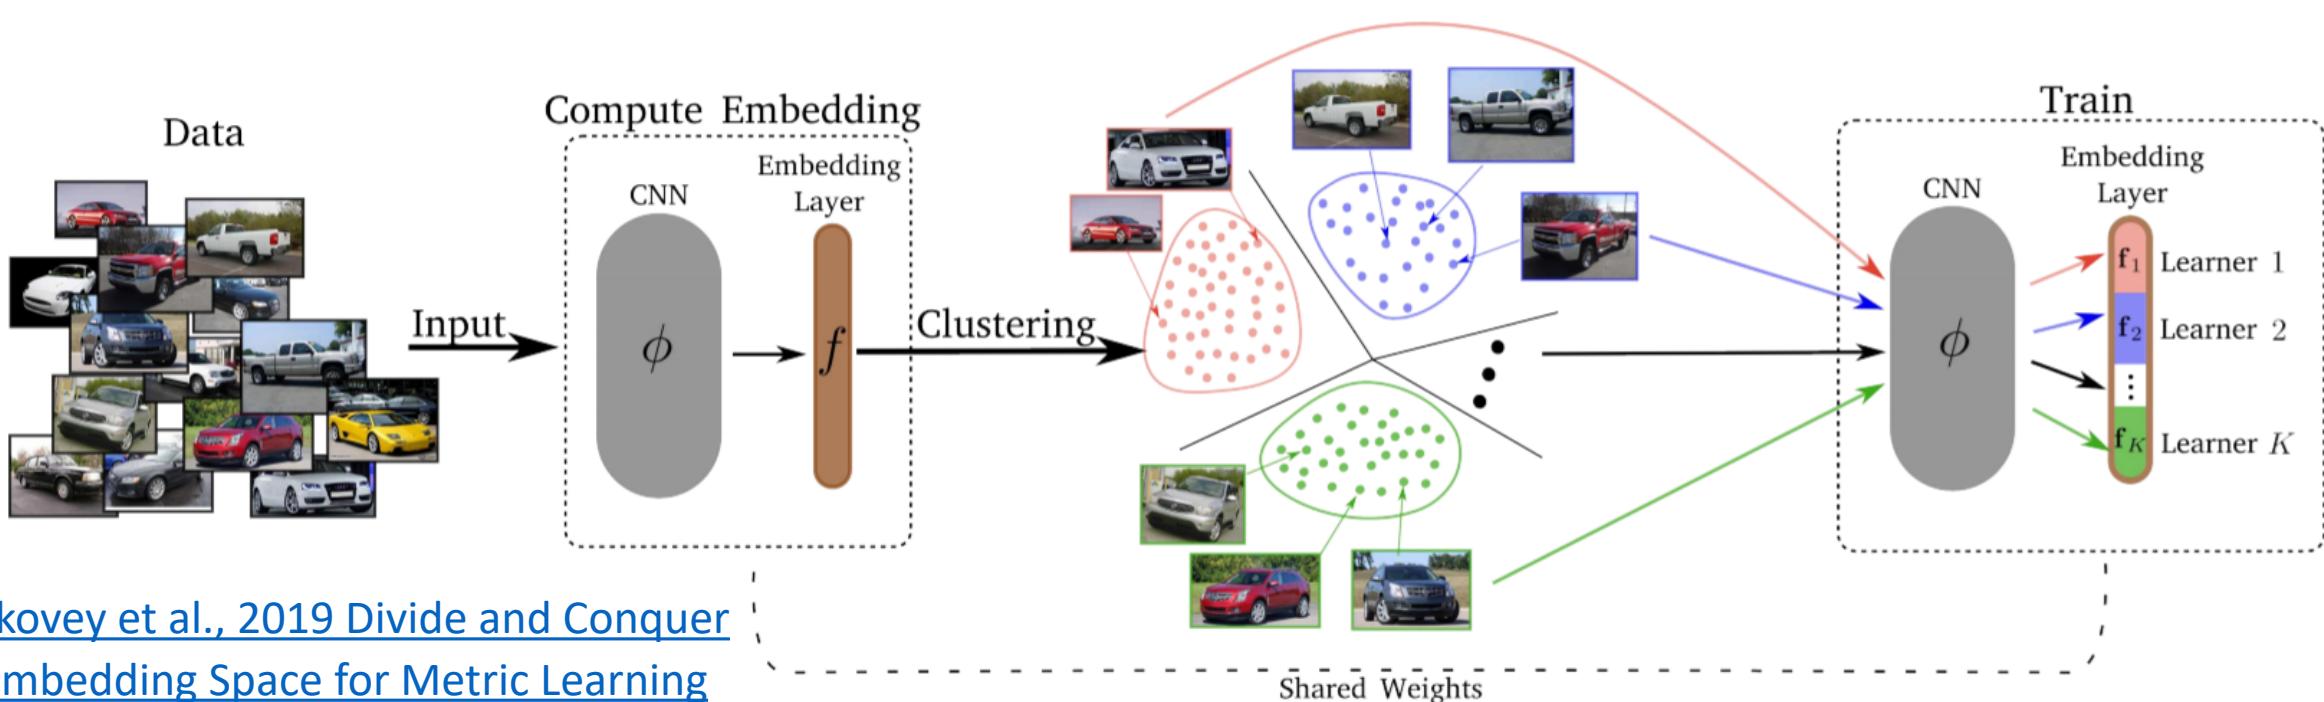


# Divide & Conquer

## DIVIDE

3) Train alternately the  $K$  learners with samples from the  $k$ -th cluster.

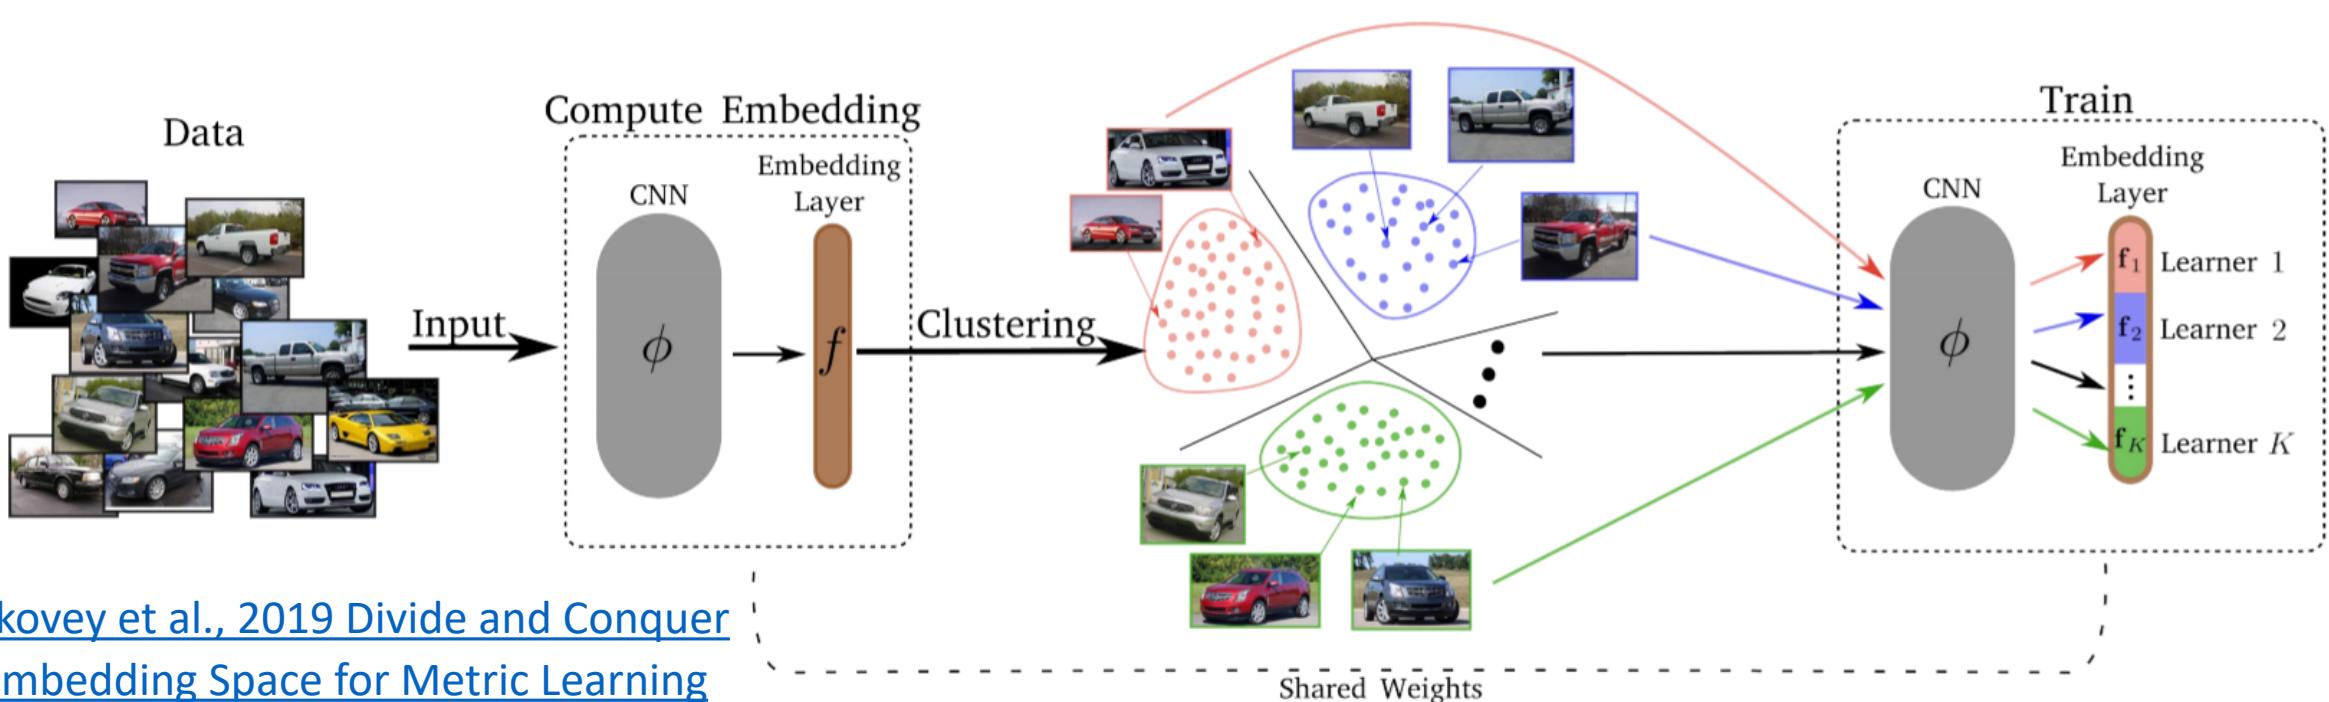
Note that each backward pass will update the parameters of the shared representation  $\phi$  and the parameters of the current learner  $f_k$ .



# Divide & Conquer

DIVIDE

- 4) Every  $T$  epochs re-cluster using the full embedding's

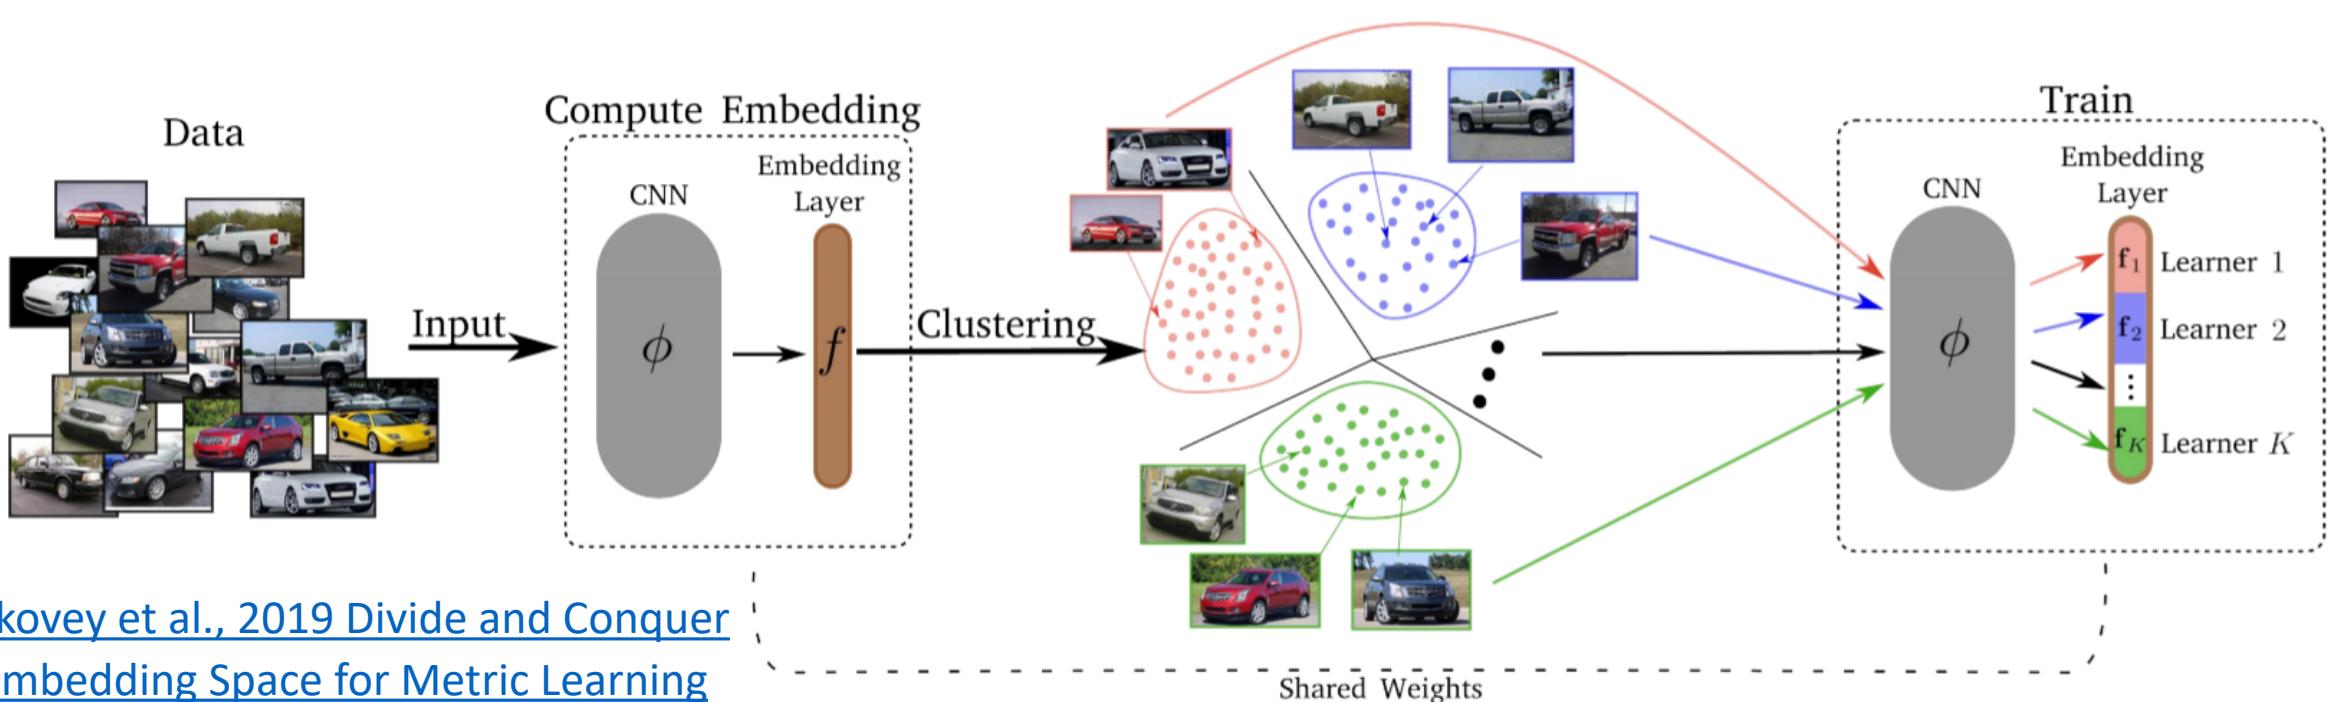


[Sanakovey et al., 2019 Divide and Conquer  
the Embedding Space for Metric Learning](#)

# Divide & Conquer

## CONQUER

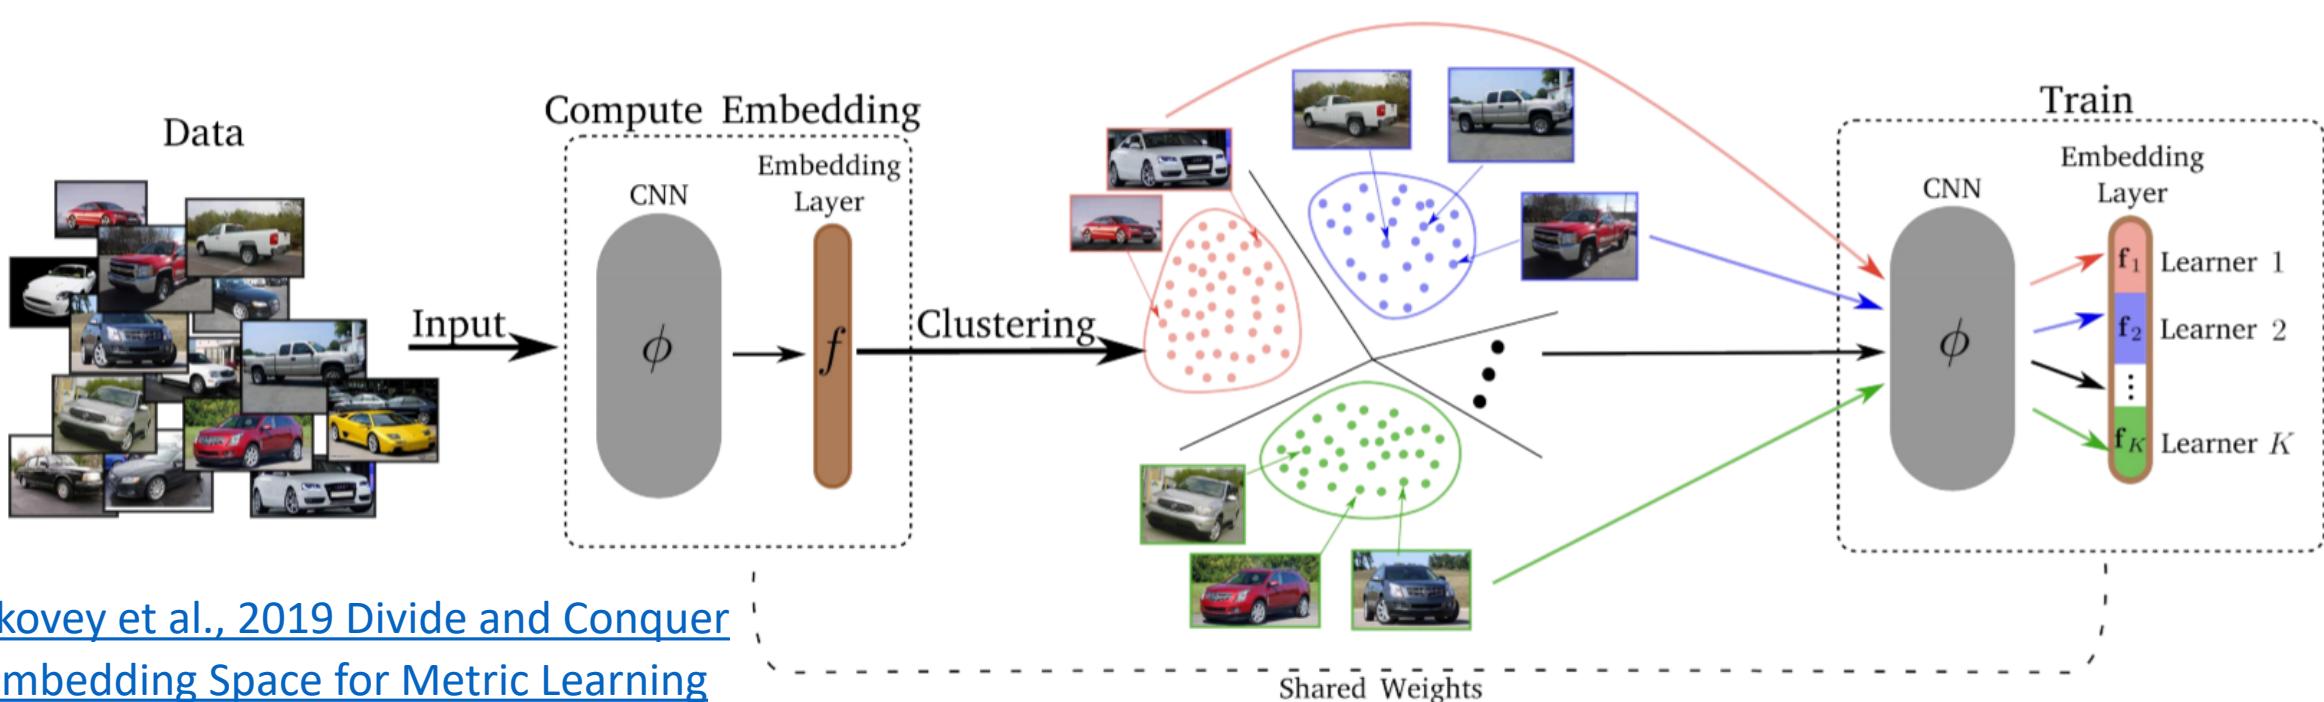
5) After convergence, merge their solution to get the full embedding space back



# Divide & Conquer

## CONQUER

- 6) Fine tune the embedding layer on the entire data set to achieve consistency between embedding of the individual learners.



# Divide & Conquer Performance

**Example:**

Results on the *n* *Stanford Online Product* [dataset](#)



# Divide & Conquer Performance

**Example:**

Results on the *n* *Stanford Online Product dataset*

HTL →  
Divide & Conquer →

R@k	1	2	4	8	NMI
Triplet Semihard [42]	51.5	63.8	73.5	82.4	53.4
LiftedStruct [32]	48.3	61.1	71.8	81.1	55.1
FacilityLoc [42]	58.1	70.6	80.3	87.8	59.0
SmartMining [14]	64.7	76.2	84.2	90.2	-
N-pairs [41]	71.1	79.7	86.5	91.6	64.0
Angular [47]	71.4	81.4	87.5	92.1	63.2
ProxyNCA [30]	73.2	82.4	86.4	88.7	64.9
HDC [51]	73.7	83.2	89.5	93.8	-
DAML (N-pairs) [6]	75.1	83.8	89.7	93.5	66.0
HTG [53]	76.5	84.7	90.4	94	-
BIER [33]	78.0	85.8	91.1	95.1	-
HTL [10]	81.4	88.0	92.7	95.7	-
DVML [28]	82.0	88.4	93.3	96.3	67.6
A-BIER [34]	82.0	89.0	93.2	96.1	-
Margin baseline [49]	79.6	86.5	91.9	95.1	69.1
<b>Ours (Margin)</b>	<b>84.6</b>	<b>90.7</b>	<b>94.1</b>	<b>96.5</b>	<b>70.3</b>
DREML [50]	86.0	91.7	95.0	97.2	76.4

**Table 2:** Recall@k for  $k = 1, 2, 4, 8$  and NMI on CARS196

# References for Ensembles

- Opitz et al., 2017, Boosting Independent Embeddings Robustly
- Elezi et al., 2020, The Group Loss for Metric Learning
- Yuan et al., 2017, Hard-Aware Deeply Cascaded Embedding
- Wang et al., 2019, Ranked list loss for deep metric learning

# Improving Similarity Learning

- Loss:
  - Contrastive vs. triplet loss vs. WDMC
- Sampling
  - Sample the space wisely diversity of classes & hard cases
- Ensembles
  - Multiple networks, each trained with a subset of triplets
- Use the Classification loss for similarity learning

# References for Classification loss

- [Movshovitz-Attias et al., 2017, No Fuss Distance Metric Learning using Proxies](#)
- [Teh et al., 2020, ProxyNCA++L Revisiting and Revitalizing Proxy Neighborhood Component Analysis](#)
- [Qian et al, 2019, SoftTriple Loss: Deep Metric Learning Without Triplet Sampling](#)
- [Elezi et al., 2020, The Group Loss for Deep Metric Learning.](#)

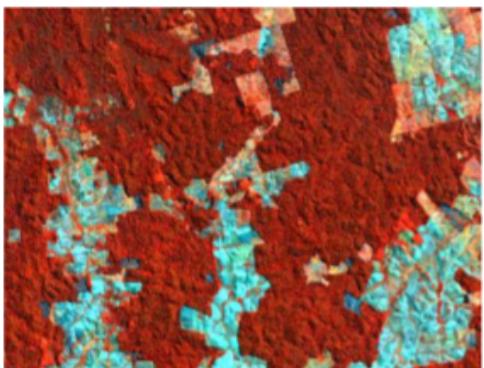
# Content

- One-shot Learning
- Example: face recognition
- Similarity Learning
- Siamese Networks
- Improving Similarity Learning
- Change Detection

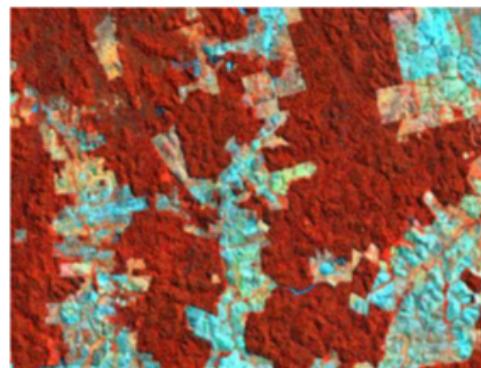
# Change Detection

Change detection involves finding the changes in images of the same location taken at different times.

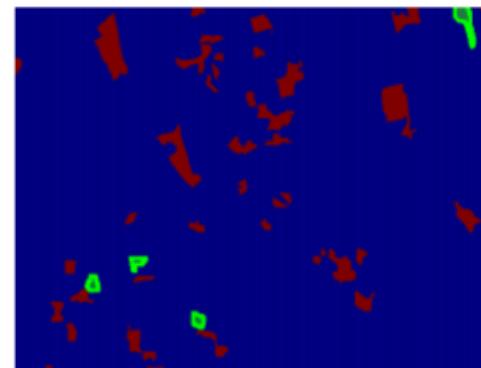
**Example:** Deforestation in Brazilian Amazon and Cerrado Biomes\*



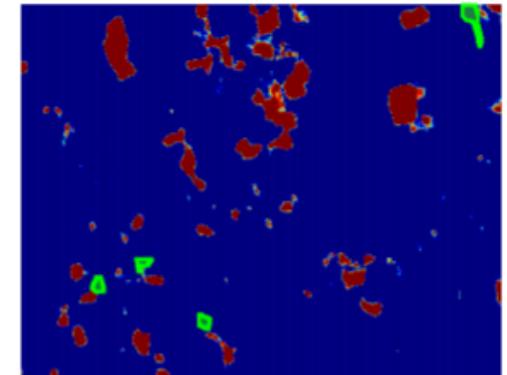
T1 (2016)



T2 (2017)



Reference



SN-4T

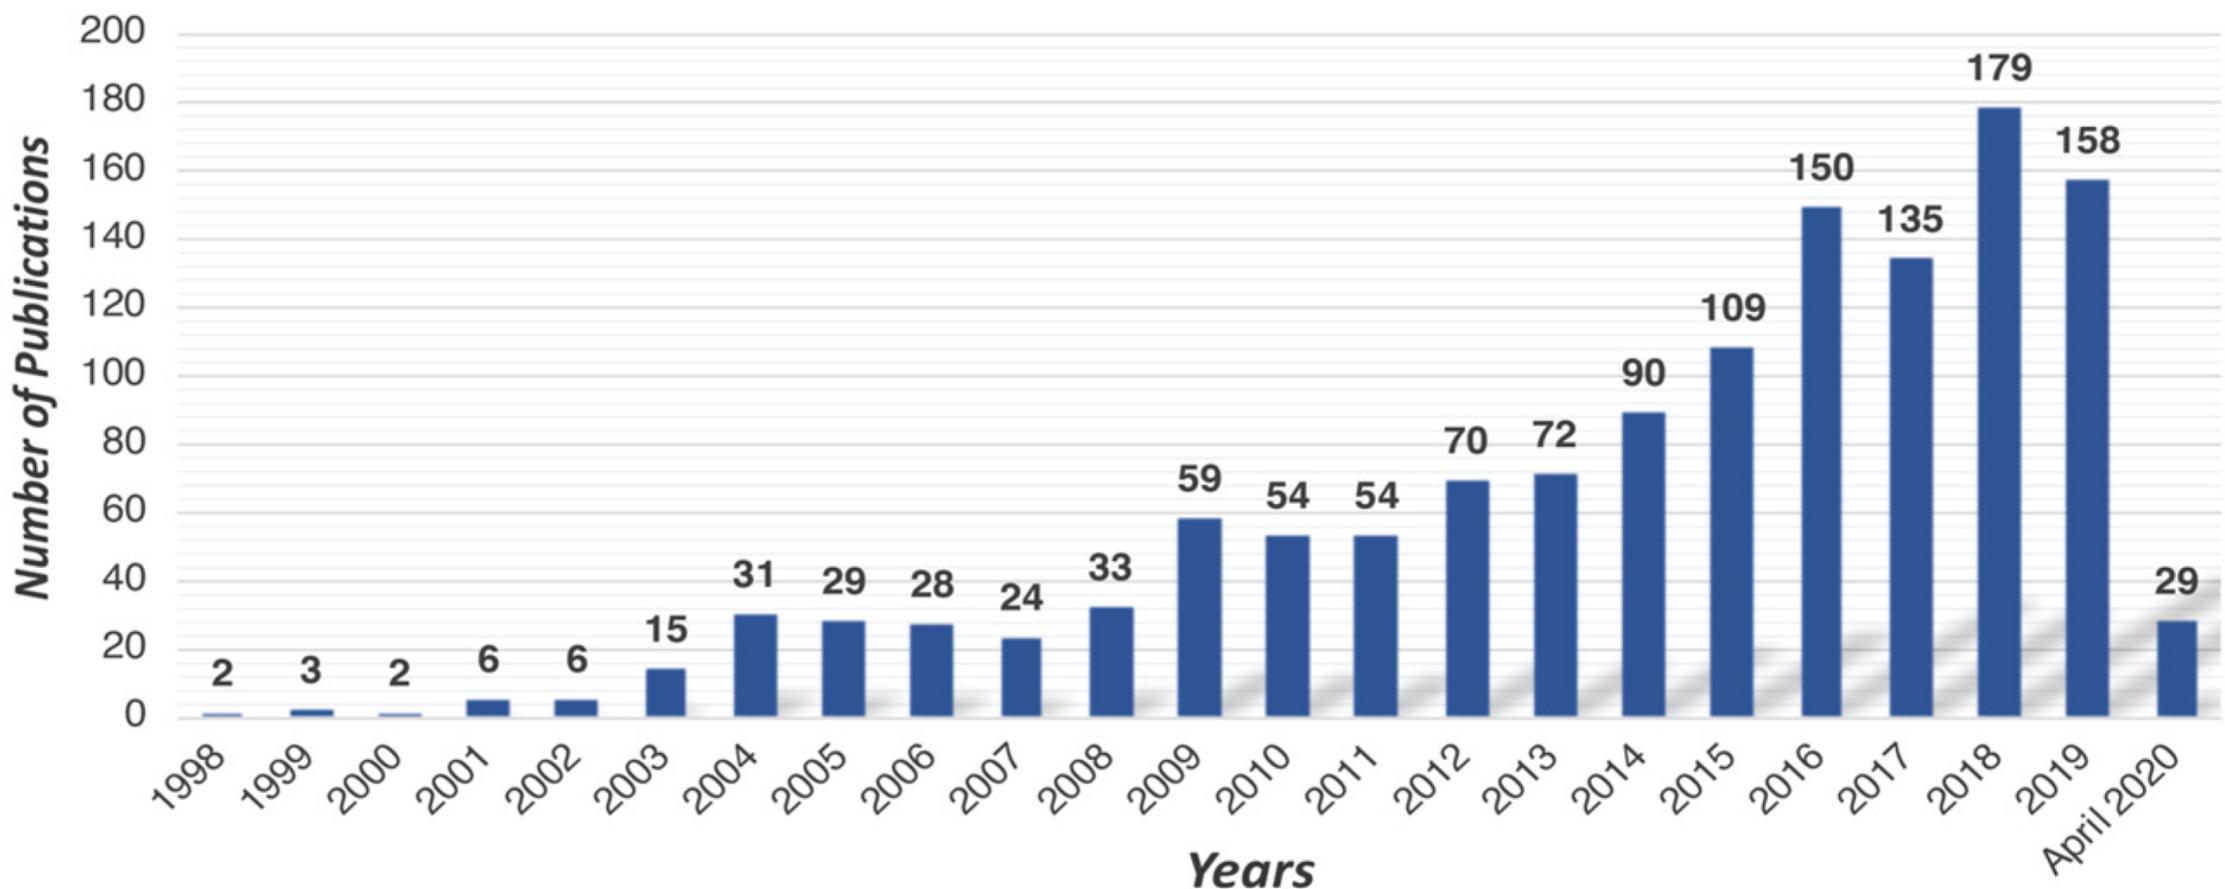
\*Source: [Adarme et al. 2020, "Evaluation of Deep Learning Techniques for Deforestation Detection in the Brazilian Amazon and Cerrado Biomes From Remote Sensing Imagery"](#)

# Change Detection Applications

- land use and land cover change (Liu et al. [2018a](#), [b](#), [c](#), [d](#), [e](#)) (Amici et al. [2017](#)),
- deforestation (Adarme et al., [2020](#)),
- urban settlements (Kleynhans et al. [2015](#)) ,
- changes as part of natural calamities (Höbling et al. [2015](#)) (Feizizadeh et al. [2017](#)),
- ...

# Research on Change Detection

Published literature on urban change detection, according to the keywords *remote sensing* and *urban change detection*.

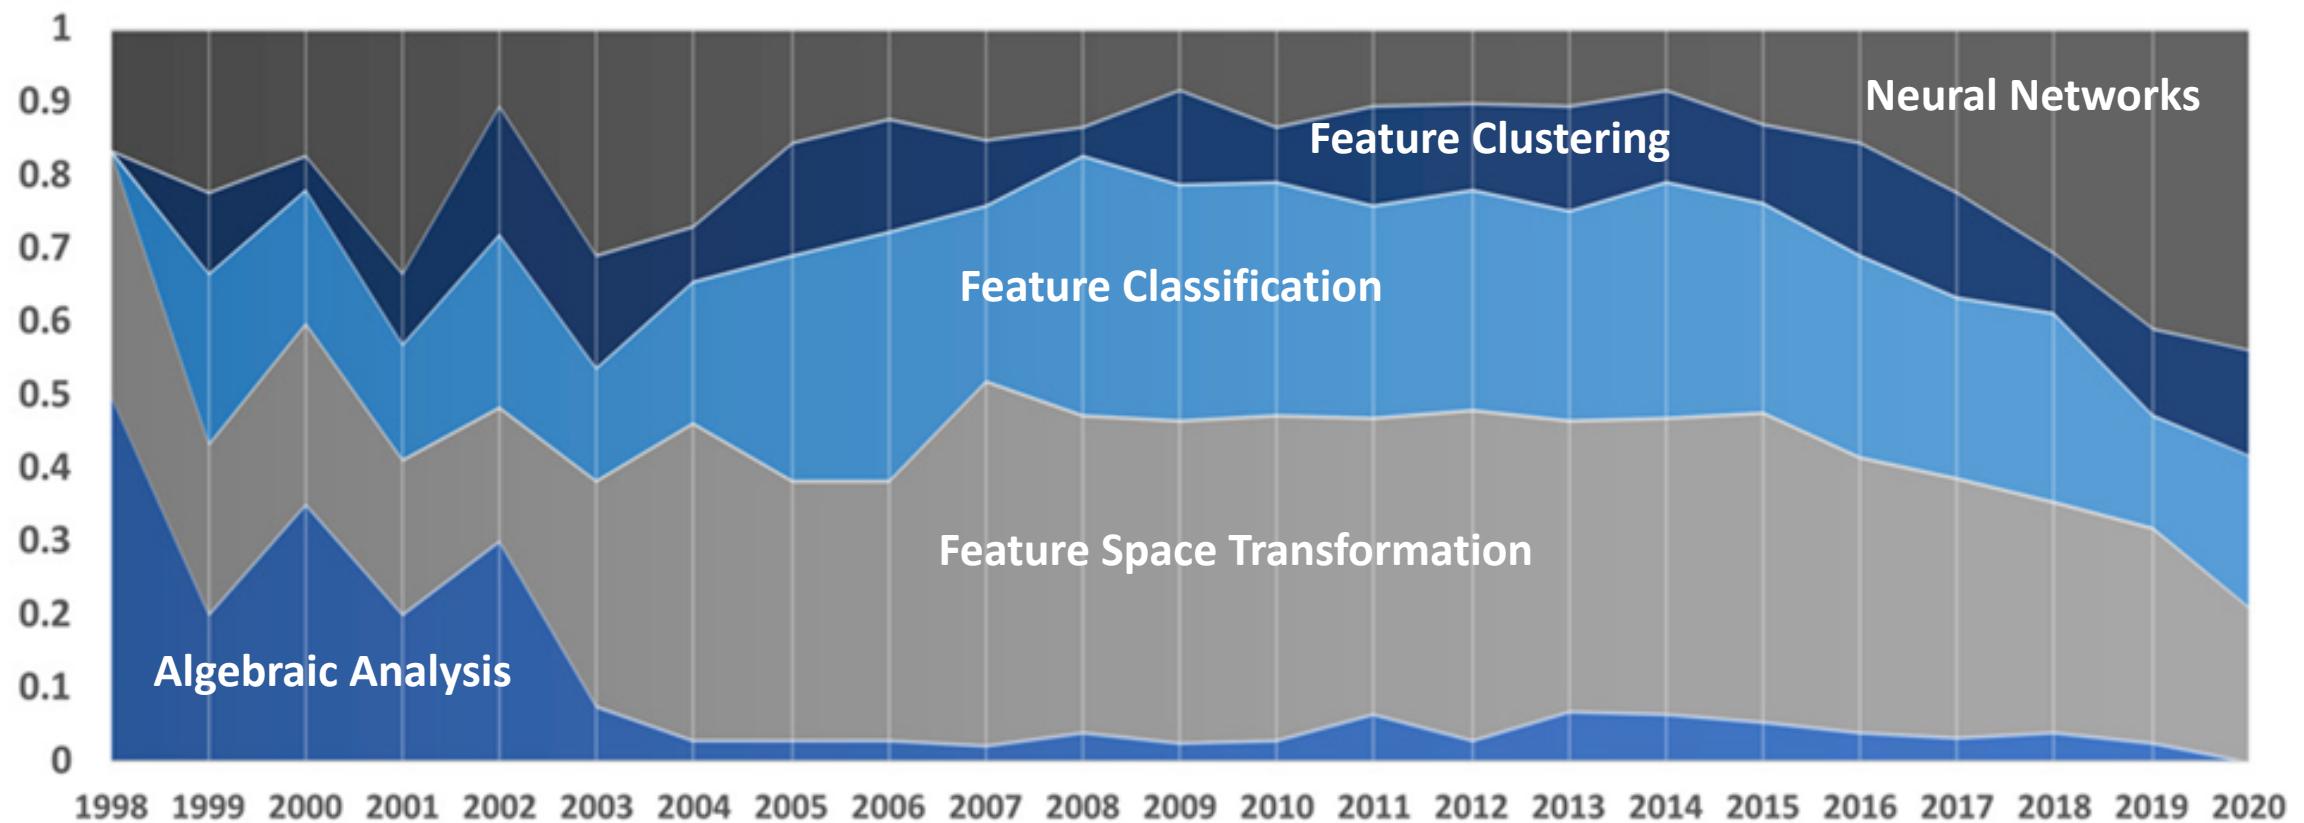


# Change Detection Databases

#	Database	Weblink	Resolution	No: of images
1.	Landsat 5 and Landsat 8 database (Zanchetta et al. <a href="#">2016</a> ), (Kant and Singh <a href="#">2017</a> )	<a href="http://glovis.usgs.gov/">http://glovis.usgs.gov/</a>	2634 × 3126	42,992
2.	PALSAR mosaic Synthetic Aperture Radar (SAR), Sentinel C band SAR and optical data (Johnson et al. <a href="#">2017</a> )	<a href="http://www.eorc.jaxa.jp/ALOS/en/palsar_fnf/fnf_index.htm">http://www.eorc.jaxa.jp/ALOS/en/palsar_fnf/fnf_index.htm</a> , <a href="https://scihub.copernicus.eu/">https://scihub.copernicus.eu/</a>	4500 × 4500	108
3.	MODIS–VCF(MOD44B), Landsat dataset (Amarnath et al. <a href="#">2017</a> )	<a href="http://glcf.umiacs.umd.edu/data/vcf/">http://glcf.umiacs.umd.edu/data/vcf/</a>	6406 × 5521	500
4.	Landsat database (Haque and Basak <a href="#">2017</a> ), (Uchenna et al. <a href="#">2017</a> )	<a href="http://www.earthexplorer.usgs.gov">www.earthexplorer.usgs.gov</a>	938 × 528	Daily image updation
5.	Landsat database (Kaliraj et al. <a href="#">2017</a> )	<a href="http://glcfapp.glc.umd.edu:8080/esdi/">http://glcfapp.glc.umd.edu:8080/esdi/</a>	6406 × 5521	800
6.	Moderate-resolution Imaging Spectroradiometer (MODIS) surface reflectance products (MOD09A1) (Qiu et al. <a href="#">2017</a> )	<a href="https://ladsweb.nascom.nasa.gov/">https://ladsweb.nascom.nasa.gov/</a>	600 × 483	Daily image updation
7.	Landsat surface reflectance images (Ye et al. <a href="#">2018</a> )	<a href="http://earth.google.com/">http://earth.google.com/</a>	4096 × 4096	Daily image updation
8.	Pleiades imagery (Suresh and Lal <a href="#">2017a</a> )	<a href="http://www.satpalda.com/">http://www.satpalda.com/</a> gallery/pleiades-imagery/	1464 × 1143	Daily image updation
9.	IR imagery and HURDAT2 data (Pradhan et al. <a href="#">2018</a> )	<a href="http://www.nrlmry.navy.mil">http://www.nrlmry.navy.mil</a> and <a href="http://www.nhc.noaa.gov/data/#hurdat">http://www.nhc.noaa.gov/data/#hurdat</a>	625 × 625	200
10.	Airborne Visible/Infrared Imaging Spectrometer (AVIRIS) sensor data (R et al. <a href="#">2017</a> )	<a href="http://aviris.jpl.nasa.gov/">http://aviris.jpl.nasa.gov/</a>	677 × 512	4000
11.	World-View2 Multispectral and Panchromatic imagery (Luo et al. <a href="#">2016</a> )	<a href="https://www.digitalglobe.com/">https://www.digitalglobe.com/</a>	1024 × 614	5000
12.	Global land cover dataset (Wang et al. <a href="#">2018a, b</a> )	<a href="http://www.globallandcover.com/">http://www.globallandcover.com/</a>	700 × 400	853
13.	Landsat images (Pandey and Khare <a href="#">2017</a> )	<a href="http://bhuvan.nrsc.gov.in">http://bhuvan.nrsc.gov.in</a>	931 × 644	Daily updation

# Methods for Change Detection

Deep Learning-based approaches are gaining prominence.



# Interesting Recent Publications

- [You, Cao and Zhou, 2020, A Survey of Change Detection Methods Based on Remote Sensing Images for Multi-Source and Multi-Objective Scenarios](#)
- [Asokan & Jude, 2019, Change detection techniques for remote sensing applications: a survey, Earth Science Informatics,](#)
- [Xudong, et al., 2019, A Conditional Adversarial Network for Change Detection in Heterogeneous Image, IEEE Geoscience and Remote Sensing Letter](#)
- [Chen et al., 2020, DASNet: Dual attentive fully convolutional siamese networks for change detection of high resolution satellite images](#)
- [Touati, Mignotte, and Dahmane, 2020, Partly Uncoupled Siamese Model for Change Detection from Heterogeneous Remote Sensing Imagery, Journal of Remote Sensing & GIS](#)
- [Zhang et al., 2019, Triplet-Based Semantic Relation Learning for Aerial Remote Sensing Image Change Detection, IEEE Geoscience and Remote Sensing Letters,](#)

# Next Lecture

**Tuesday  
Lab + 4th Programming Assignment**

*Similarity Learning*

See you next class!

